



**CME 213**  
**SPRING 2017**

Eric Darve

1891

**Stanford University**

# **SYLLABUS**

## PEOPLE

### Instructors:

- Eric Darve, ME, ICME
- Colfax and NVIDIA engineers

### Teaching assistants:

- Qing Wang, wangqing@stanford.edu
- Ruoxi Wang, ruoxi@stanford.edu

## WEB SITES

- **piazza:**
  - Forum discussion, Q&A
  - You need to register
- **canvas.stanford.edu**
  - Homework upload
  - Lecture notes, reading material
  - Computer code, etc

## **GRADING, HOMEWORK, PROJECT**

- **1 pre-requisite homework + 4 homework assignments: 65% of grade**
- **One final project: 35% of grade**
- **The first homework is due Wednesday April 12.**
- **Honor code:**  
**“that they will not give or receive aid in examinations; that they will not give or receive unpermitted aid in class work, in the preparation of reports, or in any other work that is to be used by the instructor as the basis of grading”**
- **You cannot copy someone's computer code.**
- **The work you submit must be your own.**

## COMPUTERS

- For the first few homeworks, you can use FarmShare or your own computer. You need a standard C/C++ compiler.

[https://web.stanford.edu/group/farmshare/cgi-bin/wiki/index.php/Main\\_Page](https://web.stanford.edu/group/farmshare/cgi-bin/wiki/index.php/Main_Page)

- You will also have access to the computer cluster Certainty.
- Certainty has GPUs which are needed for Homework 3 and after.

## HOW TO USE CERTAINTY

- You need to use **VPN** for Certainty.

**`https://uit.stanford.edu/service/vpn`**

- To login:

**`ssh darve@certainty-login.stanford.edu`**

- Information on cluster:

**`https://hpcc-intranet.stanford.edu/resources/`**

- Check “How do I...?” for useful information.

About the HPC Center

**Resources**

How do I....?

Software



## TIPS AND TRICKS

- How to log without a password?
- That's legit!

## STEP 1

- Say you want to log in from Host A / User a to Host B / User b.
- On A, generate a pair of authentication keys. Do not enter a passphrase:

```
a@A:~> ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/a/.ssh/id_rsa):
Created directory '/home/a/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/a/.ssh/id_rsa.
Your public key has been saved in /home/a/.ssh/id_rsa.pub.
The key fingerprint is:
3e:4f:05:79:3a:9f:96:7c:3b:ad:e9:58:37:bc:37:e4 a@A
```

## STEP 2

- The `.pub` file is safe to share because this is your public key.
- On B, create the directory `~/ssh`
- Append your public key created on A to `authorized_keys` on B.

```
a@A:~> cat .ssh/id_rsa.pub | ssh b@B 'cat >> .ssh/authorized_keys'
```

# RUNNING A CODE ON CERTAINTY

See

<https://hpcc-intranet.stanford.edu/resources/using-the-hpc-clusters-101/>

## Step 1: load modules with compiler and libraries

**\$ module avail**

**\$ module list**

**\$ module load intel/16**

```
[[darve@certainty-c ~]$ module avail
----- /share/apps/modules/modulefiles -----
ansys/15.0                               gnu/4.9.2
ansys/16.1                               gnu/5.2
ansys/18.0                               hdf5/1.8.13-mvapich2-2.0rc1-intel-14
boost/1.55.0-mvapich2-2.0rc1-intel-14   hdf5/1.8.17-mvapich2-2.2b-intel-16
boost/1.63.0-mvapich2-2.2b-intel-16     intel/14
clang/3.5.2                             intel/15
cmake/2.8.12.2                           intel/16
comsol/44                                maple/2016
comsol/51                                matlab/R2014a
comsol/52                                matlab/R2015b
convergecfd/2.2.0                          mvapich2/2.0rc1-intel-14
cuda/7.0.28                               mvapich2/2.1a-gcc-4.4.7
cuda/7.5.18                               mvapich2/2.1-intel-15
dakota/6.2.0                             mvapich2/2.1rc1-intel-15
ensight/10.1.6b                           mvapich2/2.2b-gnu-4.4.7-11
freefem++/3.30                            mvapich2/2.2b-gnu-4.7.2
freefem++/3.38-1                           mvapich2/2.2b-gnu-4.9.2
freefem++/3.46-gnu-4.9.2                  mvapich2/2.2b-gnu-5.2
freefem++/3.47-gnu-4.9.2                  mvapich2/2.2b-intel-16
freefem++/3.47-mvapich2-2.2b-gnu-5.2    null
                                          openmpi/1.8.6-intel-15
                                          paraview/4.1.0
                                          parmetis/4.0.3-mvapich2-2.0rc1-intel-14
                                          parmetis/4.0.3-mvapich2-2.1rc1-intel-15
                                          parmetis/4.0.3-mvapich2-2.2b-intel-16
                                          petsc/3.0.0-p12-mvapich2-2.0rc1-intel-14
                                          petsc/3.0.0-p12-mvapich2-2.0rc1-intel-14-DEBUG
                                          petsc/3.0.0-p12-mvapich2-2.1rc1-intel-15
                                          petsc/3.4.4-mvapich2-2.0rc1-intel-14
                                          petsc/3.4.4-mvapich2-2.0rc1-intel-14-DEBUG
                                          petsc/3.7.2-mvapich2-2.2b-gnu-4.9.2
                                          petsc/3.7.2-mvapich2-2.2b-gnu-5.2
                                          pointwise/18.0R1
                                          python/2.7.8
                                          scons/2.5.0
                                          simvascular/2016-04-06
                                          tecplot/14.0.2
                                          tecplot/14.1.0.51525
                                          tecplot/2016-R2-EX
                                          totalview/2016.01.0
                                          totalview/2016.01.0
```

## RUNNING A CODE ON CERTAINTY

### Step 2:

- Compile your code
- Submit a job using `qsub` to the GPU queue:

```
$ qsub -q gpu submit.sh
```

```
$ qstat -q gpu
```

```
$ qstat -u darve
```

## EXAMPLE OF SCRIPT: SUBMIT.SH

```
#!/bin/bash

#PBS -N cme213
#PBS -e cme213.err
#PBS -o cme213.out
#PBS -l nodes=1:ppn=24
#PBS -V

PBS_O_WORKDIR='/home/darve/CME213/Lecture_03/code'
export PBS_O_WORKDIR

## # -----
## # BEGINNING OF EXECUTION
## # -----


echo The master node of this job is `hostname`
echo The working directory is `echo $PBS_O_WORKDIR`
echo This job runs on the following nodes:
echo `cat $PBS_NODEFILE`
echo
echo Output from code
echo -----

## # end of information preamble

cd $PBS_O_WORKDIR
./matrix_prod -n 512 -p 24
```

## THE FINAL PROJECT

- You will be given a final project to work on towards the end of the quarter.
- Everyone will work on the same project.
- The final project will be on neural network for machine learning.
- The project will involve multi-GPU programming with CUDA and MPI.

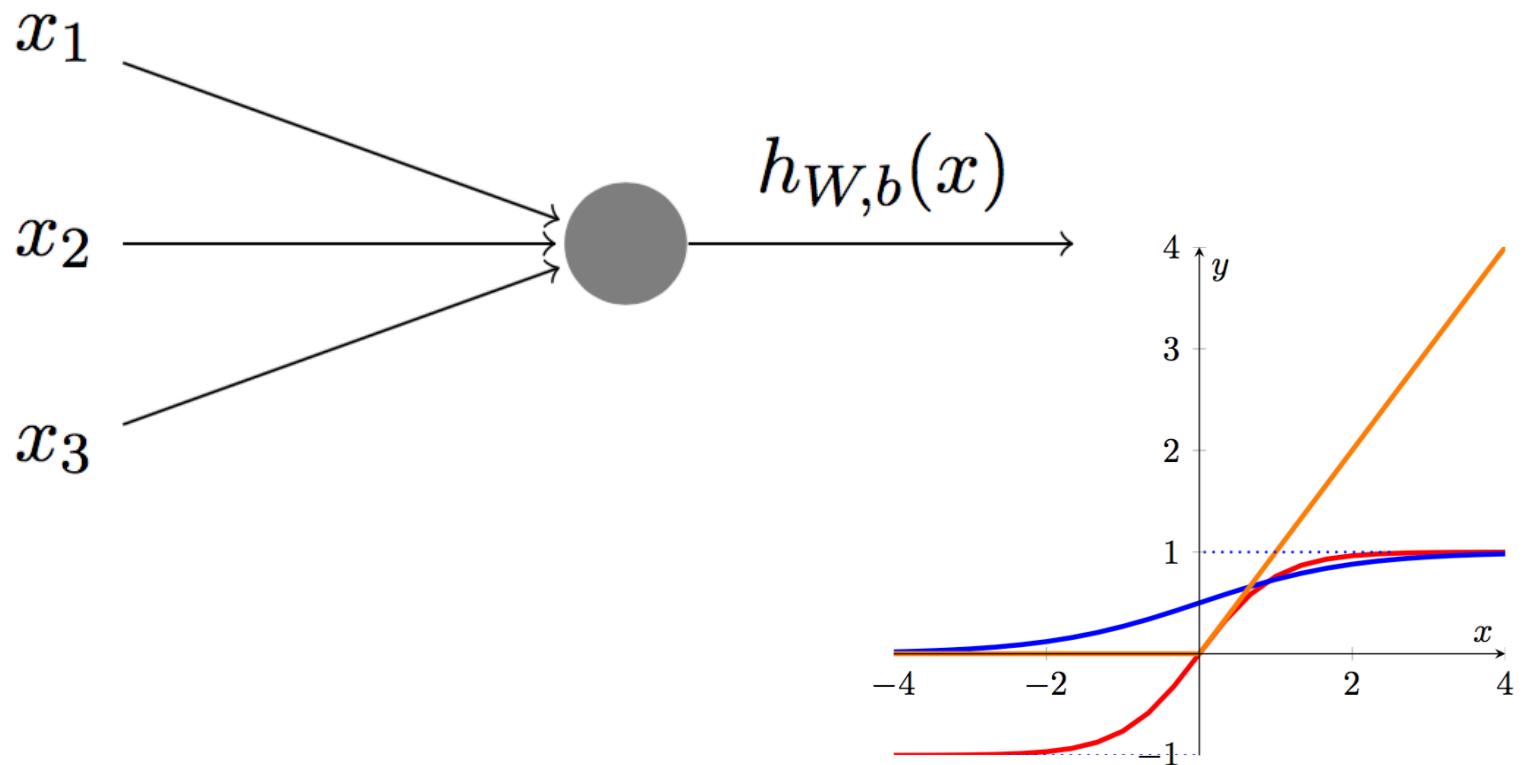
## NEURAL NETWORKS

- Simple but powerful
- “Mimic” neurons in the brain: a network of nodes that receive inputs and generate outputs.



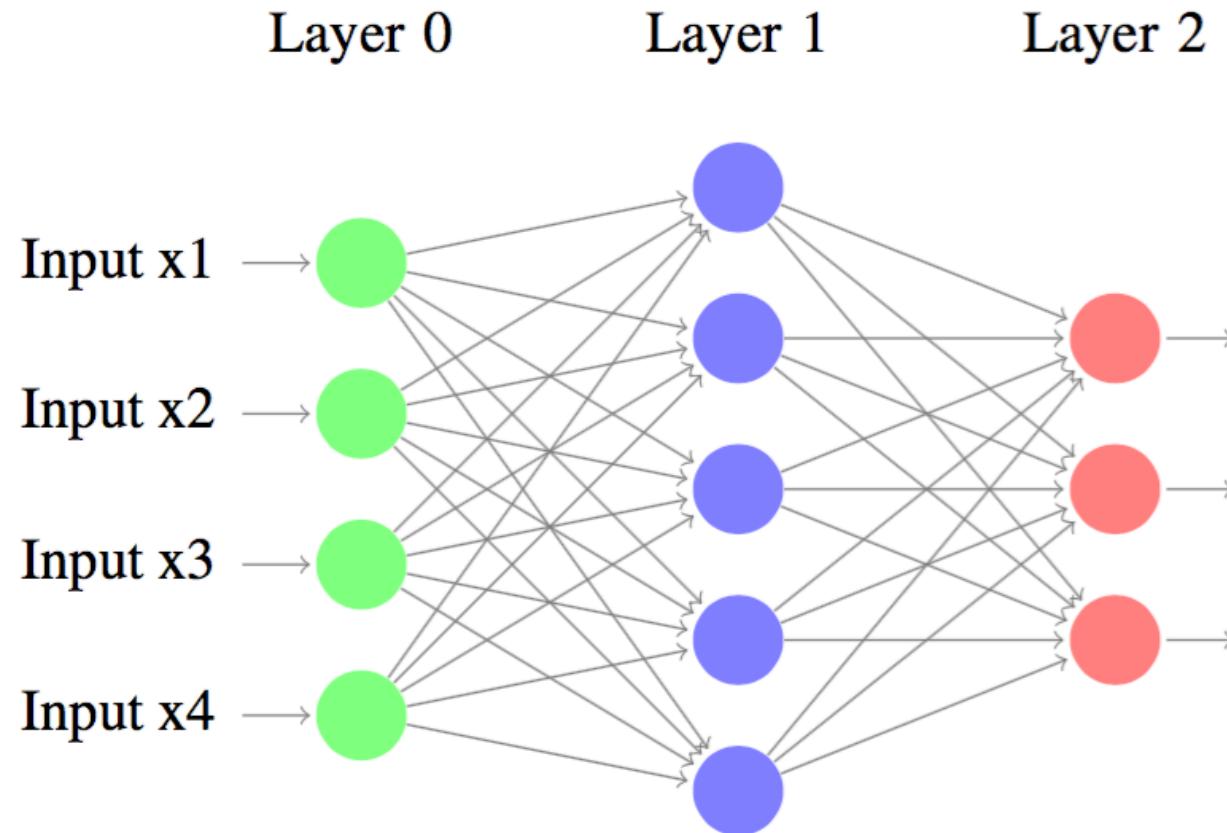
## WHAT'S A NEURON EXACTLY?

- It's a highly simplified model of a biological neuron.

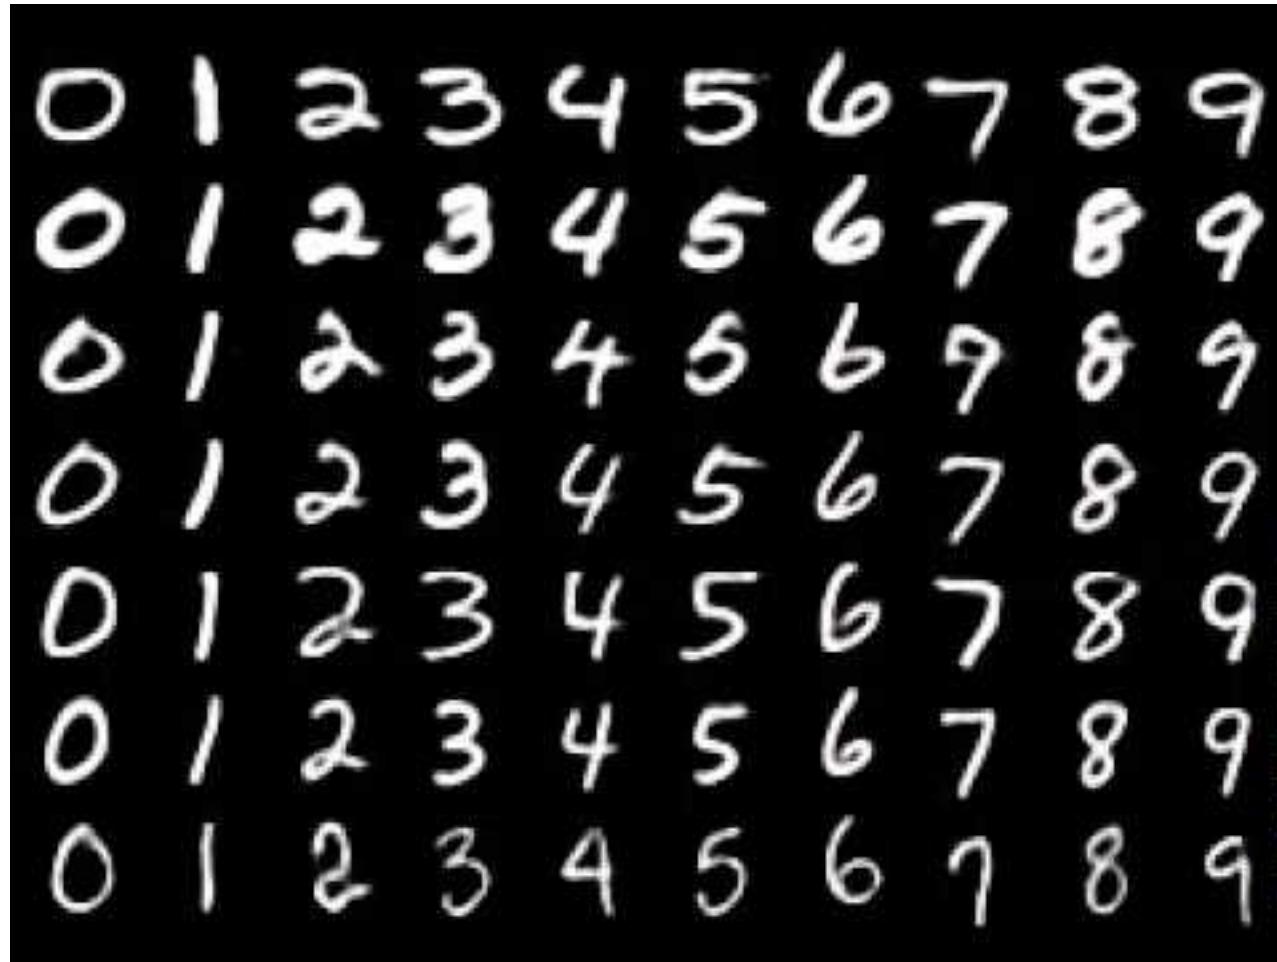


—  $\tanh(x)$  — Sigmoid — ReLU

**Assemble many neurons to create a neural network.**



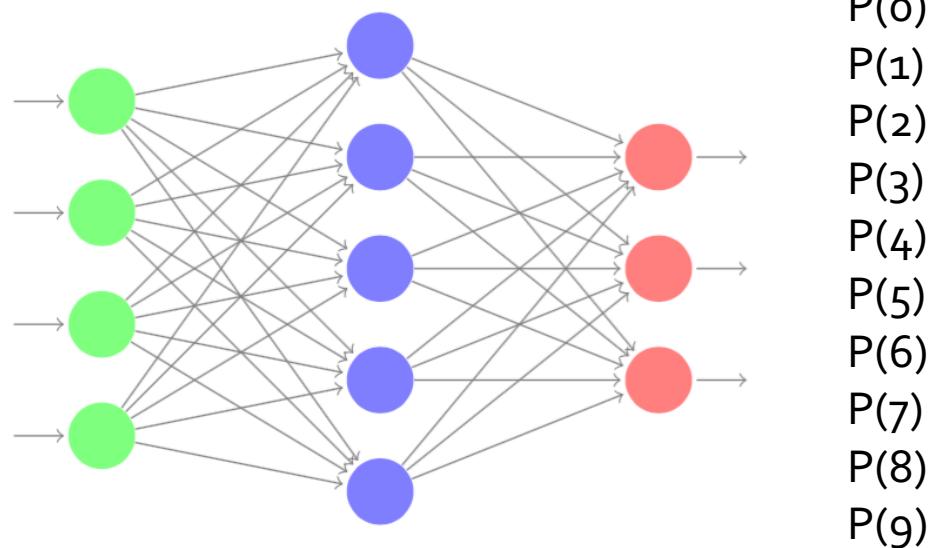
## INPUT TO NEURAL NETWORK: HANDWRITTEN DIGIT



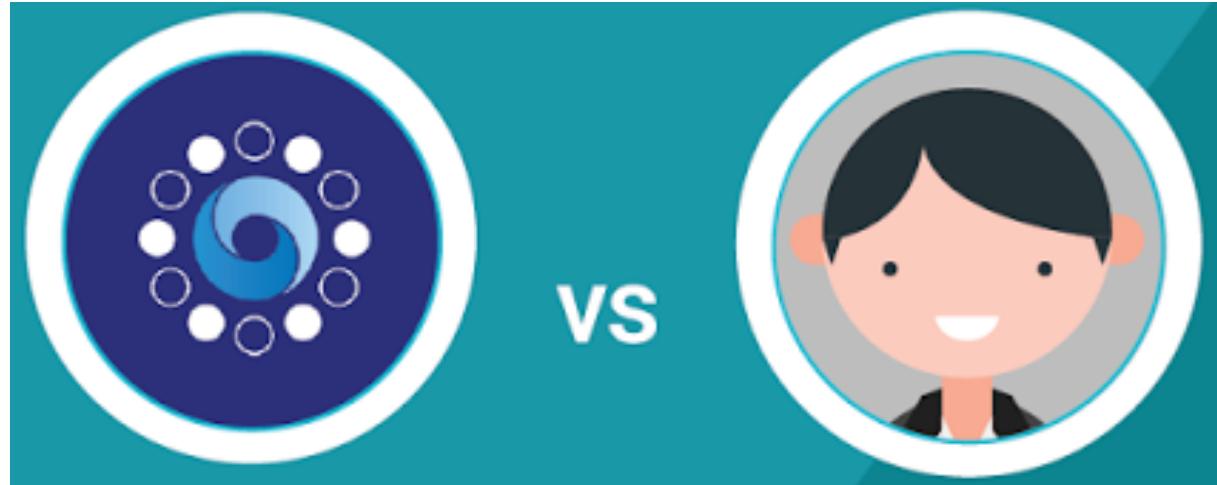
## OUTPUT: PROBABILITY OF DIGIT

- The network gives a probability vector of size 10
- Entry  $i$  is the probability that the image is showing the digit  $i$ .

0	1	2	3	4	5	6	7	8	9
0	1	<b>2</b>	4	5	6	7	8	9	
0	1	<b>2</b>	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9



## ALPHAGO

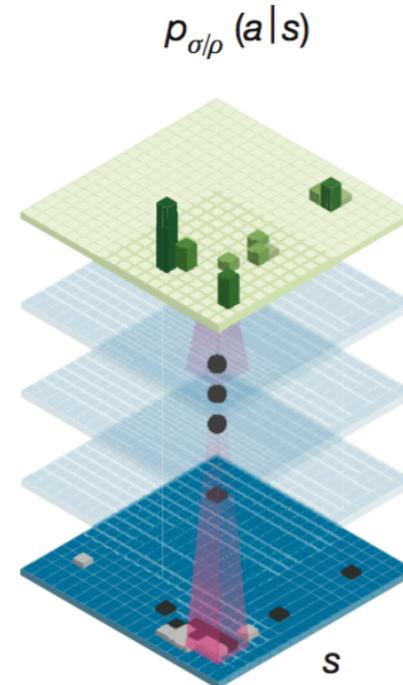


- AlphaGo (a computer program written by Google DeepMind) took on and defeated legendary Go player, Lee Sedol (9-dan professional with 18 world titles).
- The final score was 4 to 1.

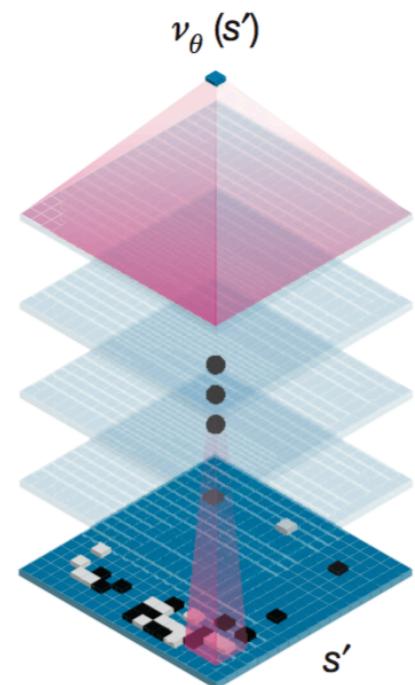
# ALPHAGO

- We will talk more about this later in the quarter.
- DeepMind used deep learning to create a program to play Go.
- It uses parallel computing and GPUs to use neural networks to play Go.
- Hardware: 1,202 CPUs and 176 GPUs.

Policy network

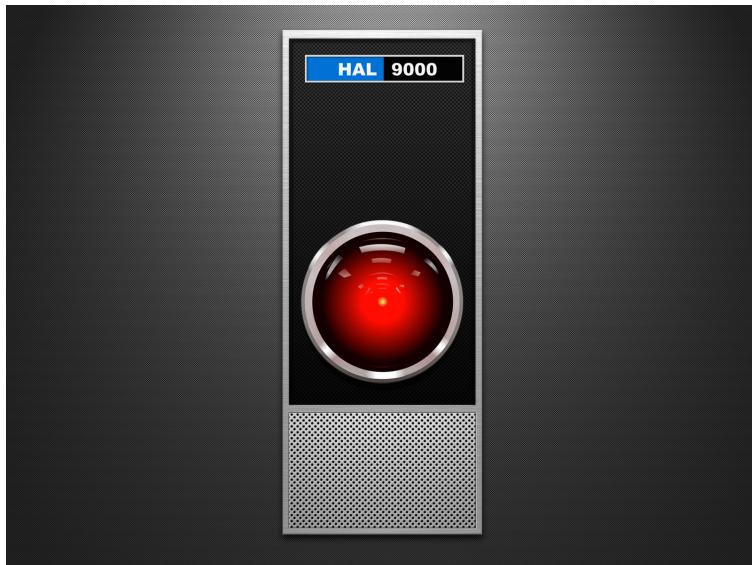


Value network



**“Although we have programmed this machine to play, we have no idea what moves it will come up with. Its moves are an emergent phenomenon from the training. We just create the data sets and the training algorithms. But the moves it then comes up with are out of our hands—and much better than we, as Go players, could come up with.”**

*DeepMind research scientist Thore Graepel*



## Books

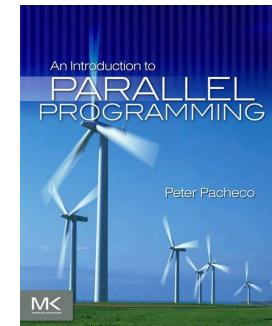
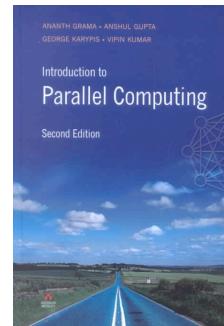
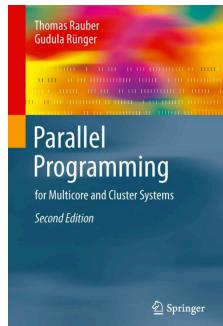
**Good news: all books are available electronically from the Stanford Library. Just go to:**

**<http://searchworks.stanford.edu/>**



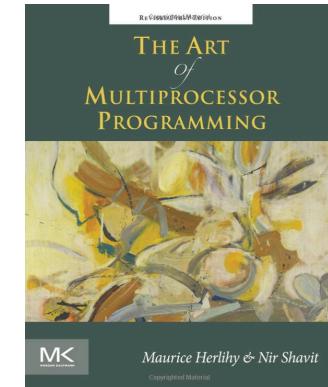
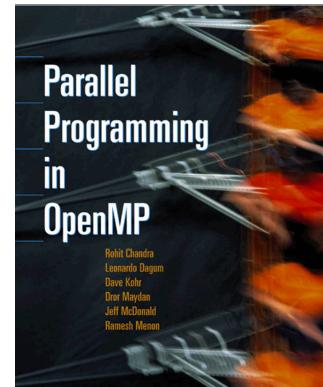
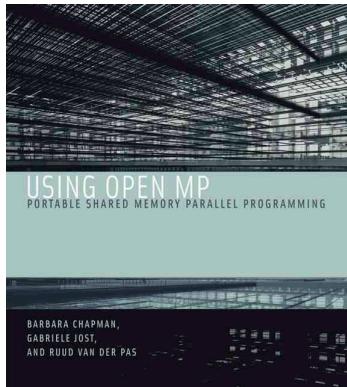
# OPENMP, MPI, PARALLEL PROGRAMMING

- *Parallel Programming for Multicore and Cluster Systems*, Rauber and Rünger. Applications focus mostly on linear algebra.
- *Introduction to Parallel Computing*, Grama, Gupta, Karypis, Kumar. Wide range of applications from sort to FFT, linear algebra and tree search.
- *An introduction to parallel programming*, Pacheco. More examples and less theoretical. Applications include n-body codes and tree search.



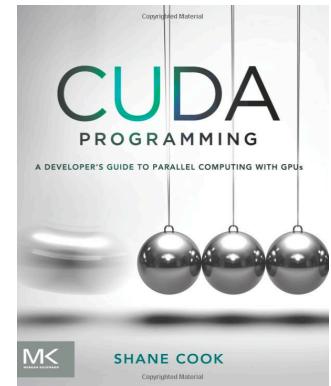
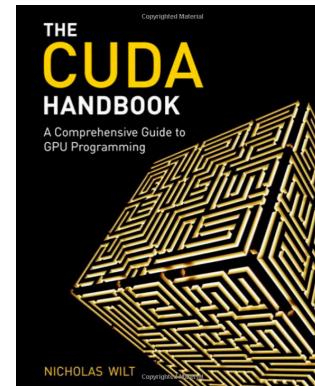
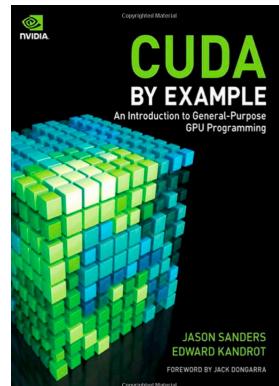
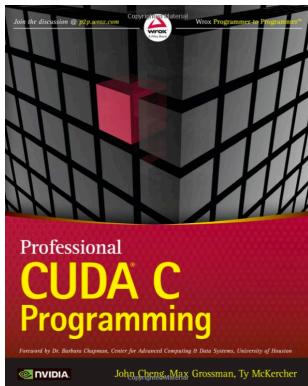
## OPENMP AND MULTICORE Books

- *Using OpenMP: portable shared memory parallel programming*, Chapman, Jost, van der Pas. Advanced coverage of OpenMP.
- *Parallel Programming in OpenMP*, Chandra, Menon, Dagum, Kohr, Maydan, McDonald; a bit outdated.
- *The art of multiprocessor programming*, Herlihy, Shavit. Specializes on advanced multicore programming.



## CUDA BOOKS

- *Professional CUDA C Programming*, Cheng, Grossman, McKercher; recent book, recommended for this class
- *CUDA by Example: An Introduction to General-Purpose GPU Programming*, Sanders, Kandrot
- *CUDA Handbook: A Comprehensive Guide to GPU Programming*, Wilt
- *CUDA Programming: A Developer's Guide to Parallel Computing with GPUs*, Cook



## CUDA BOOKS: PREFERRED REFERENCES

- CUDA online documentation:

<http://docs.nvidia.com/cuda/index.html>

<https://developer.nvidia.com/cuda-education-training>

- Uploaded reading material:

**CUDA\_C\_Best\_Practices\_Guide.pdf**

**CUDA\_C\_Programming\_Guide.pdf**

## WHAT THIS CLASS IS ABOUT

- We will focus on how to program:
  - Multicore processors, e.g., desktop processors: Pthreads, OpenMP.
  - Introduction to Xeon Phi and Knights Landing
  - NVIDIA graphics processors using CUDA.
  - Computer clusters using MPI.
- We will cover some numerical algorithms for illustration: sort, linear algebra, basic parallel primitives.

## WHAT THIS CLASS IS NOT ABOUT

- Parallel computer architecture
- Parallel design patterns and programming models
- Parallel numerical algorithms. See *CME 342: Parallel Methods in Numerical Analysis*

## WHAT THIS CLASS REQUIRES

- Some basic knowledge of UNIX (ssh, makefile, etc)
- Good knowledge of C and C++ (including pointers, templates)
- Proficiency in scientific programming, including debugging and testing

# WHY PARALLEL COMPUTING?

## WHY PARALLEL COMPUTING?

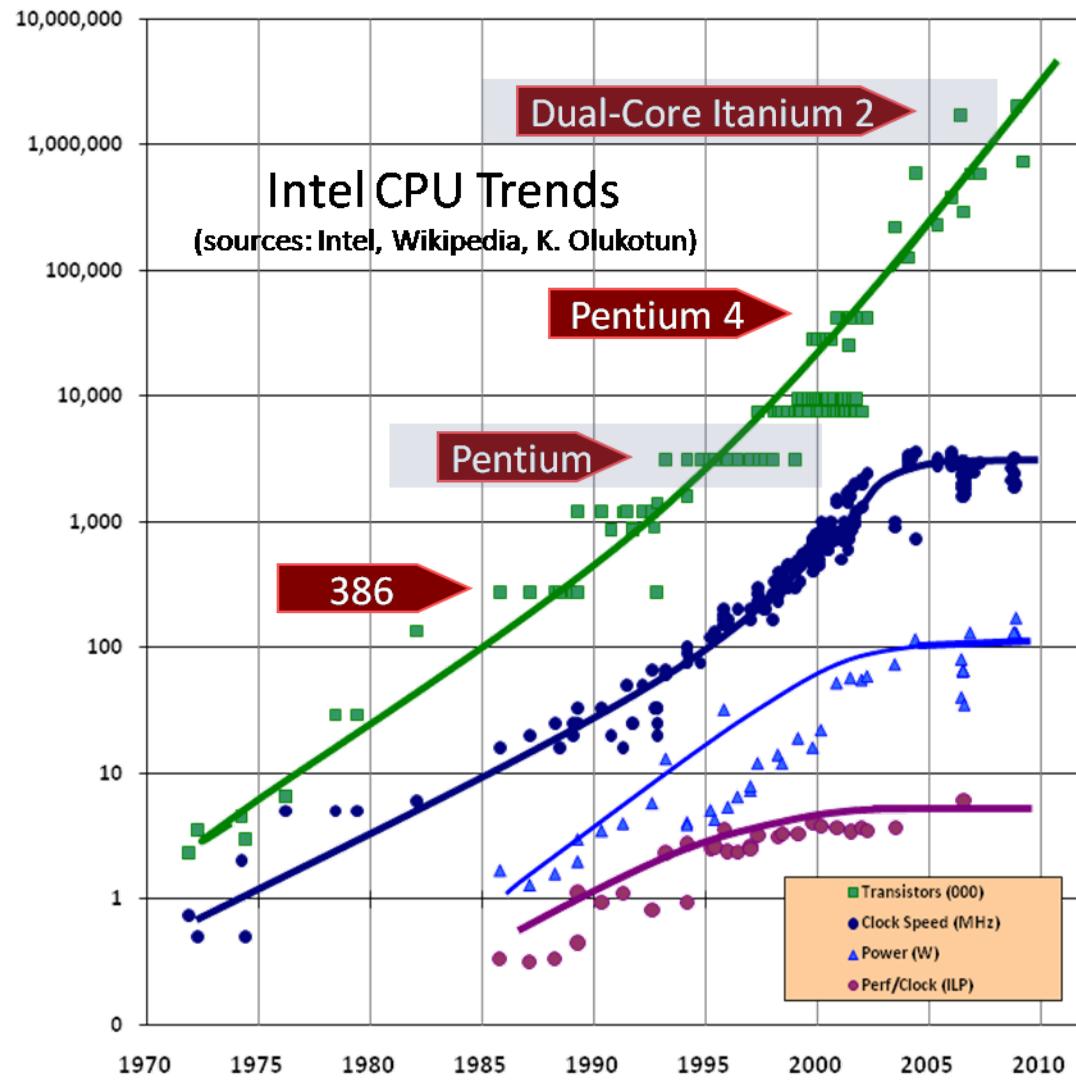
- Parallel computing has existed for a long time but until recently it was a specialized area that concerned only a small fraction of engineers.
- Nowadays, parallel computing is a dominant player in scientific and large scale computing.
- What happened?

## WHY PARALLEL COMPUTING?

- **Gordon Moore 1965:** the number of transistors on a chip shall double every 18-24 months.
- This has been valid for more than 40 years.
- This increase in the number of transistors has been accompanied by an increase in clock speed.



# INTEL MICROPROCESSOR TRENDS



## PERFORMANCE INCREASE

- Increase in transistor density is limited by:
  - Leakage current increases
  - Power consumption increases
  - Heat generated increases
- In addition, the **memory access time** has not been reduced at a rate comparable with processing speed (processor clock period).
- New ways need to be found.
- The most promising approach is to have multiple cores on a single processor.

## SEQUENTIAL VS PARALLEL



**The Sequential Giant**  
Build bigger, meaner,  
faster processor

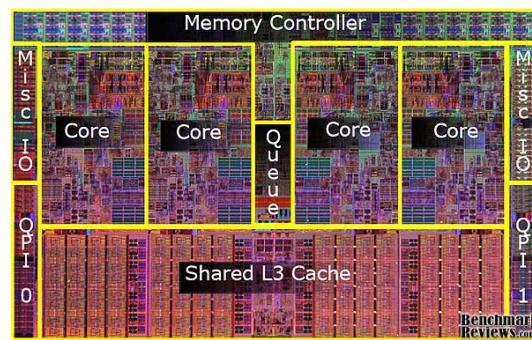
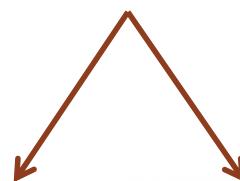


**The Parallel Ants**  
Lots of smaller, slower cores

## PARALLEL COMPUTING EVERYWHERE: MOBILE DEVICES



## LAPTOPS AND DESKTOPS

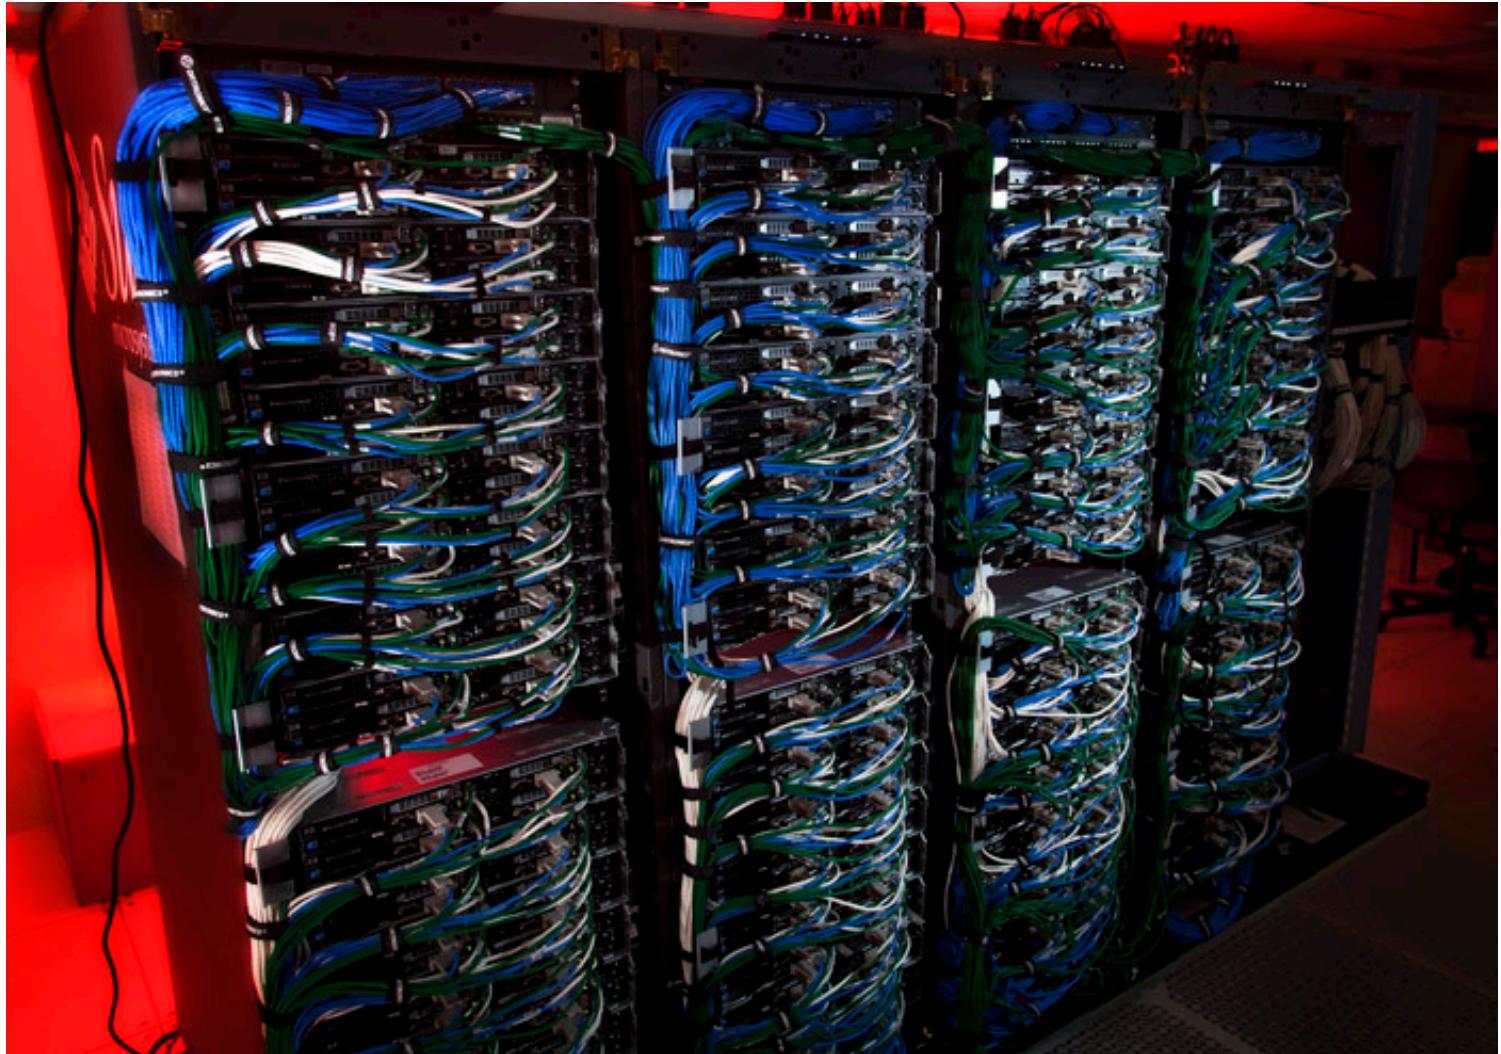


multicore

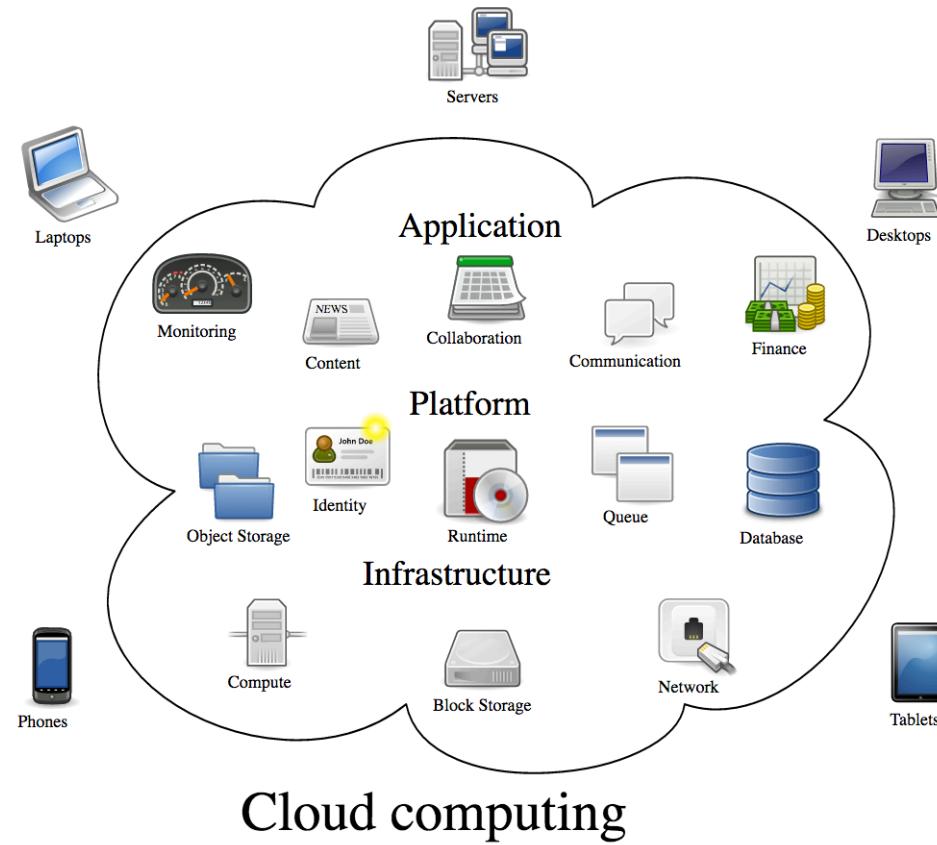


GPU

# CLUSTERS



# CLOUD COMPUTING



Cloud computing is a type of Internet-based computing that provides shared computer processing resources and data to computers and other devices on demand.

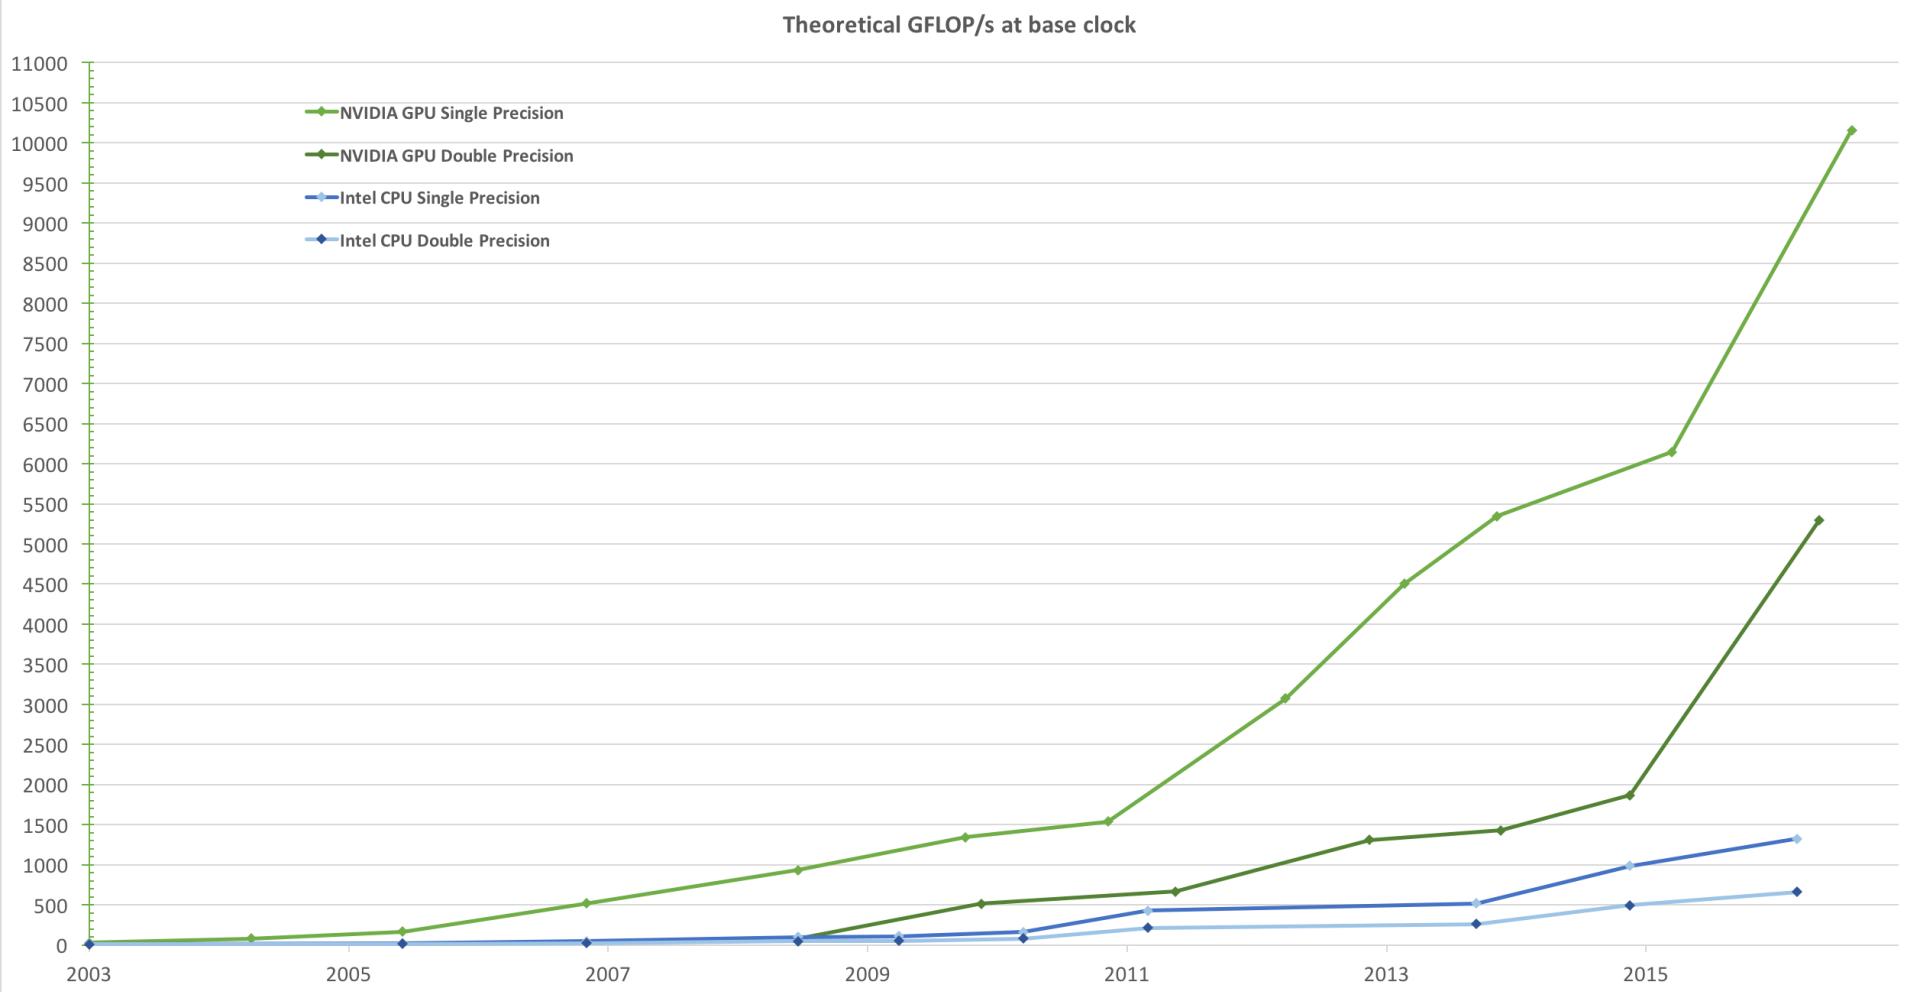
# SUPERCOMPUTERS



- Titan
- Cray Supercomputer, Oak Ridge National Lab

<b>Active</b>	Became operational October 29, 2012
<b>Sponsors</b>	US DOE and NOAA (<10%)
<b>Operators</b>	Cray Inc.
<b>Location</b>	Oak Ridge National Laboratory
<b>Architecture</b>	18,688 AMD Opteron 6274 16-core CPUs 18,688 Nvidia Tesla K20X GPUs
<b>Power</b>	8.2 MW
<b>Operating system</b>	Cray Linux Environment
<b>Space</b>	404 m <sup>2</sup> (4352 ft <sup>2</sup> )
<b>Memory</b>	693.5 TiB (584 TiB CPU and 109.5 TiB GPU)
<b>Storage</b>	40 PB, 1.4 TB/s IO Lustre filesystem
<b>Speed</b>	17.59 petaFLOPS (LINPACK) 27 petaFLOPS theoretical peak
<b>Cost</b>	\$97 million
<b>Ranking</b>	TOP500: #3, June 2016 <sup>[1]</sup>

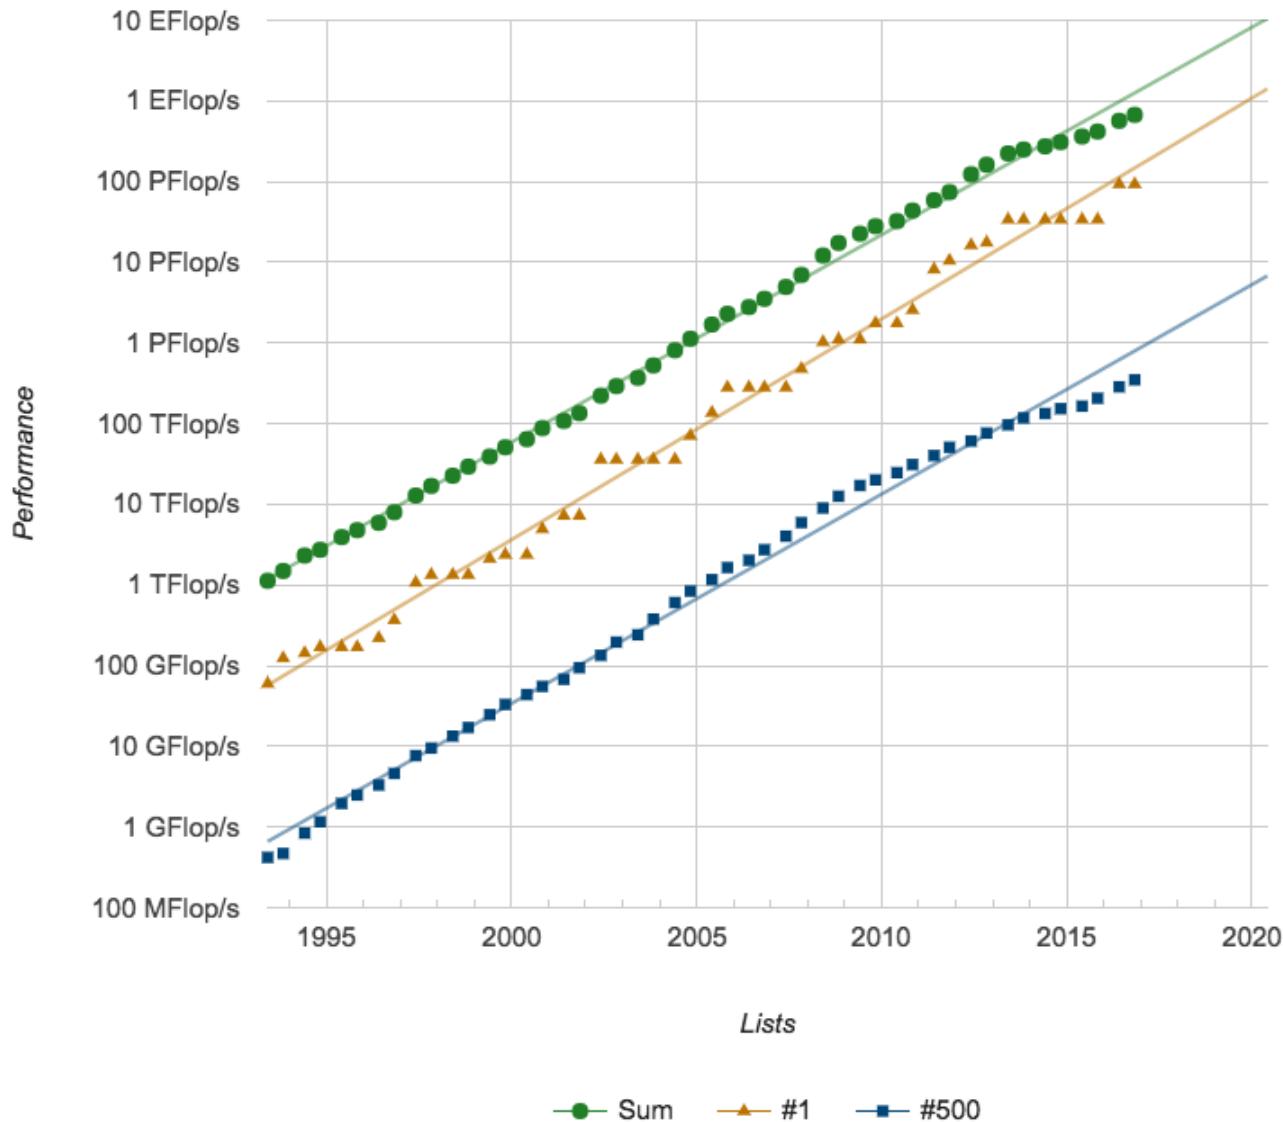
# PERFORMANCE THROUGH PARALLEL PROCESSING



# **STATISTICS ABOUT THE TOP 500 SUPERCOMPUTERS**

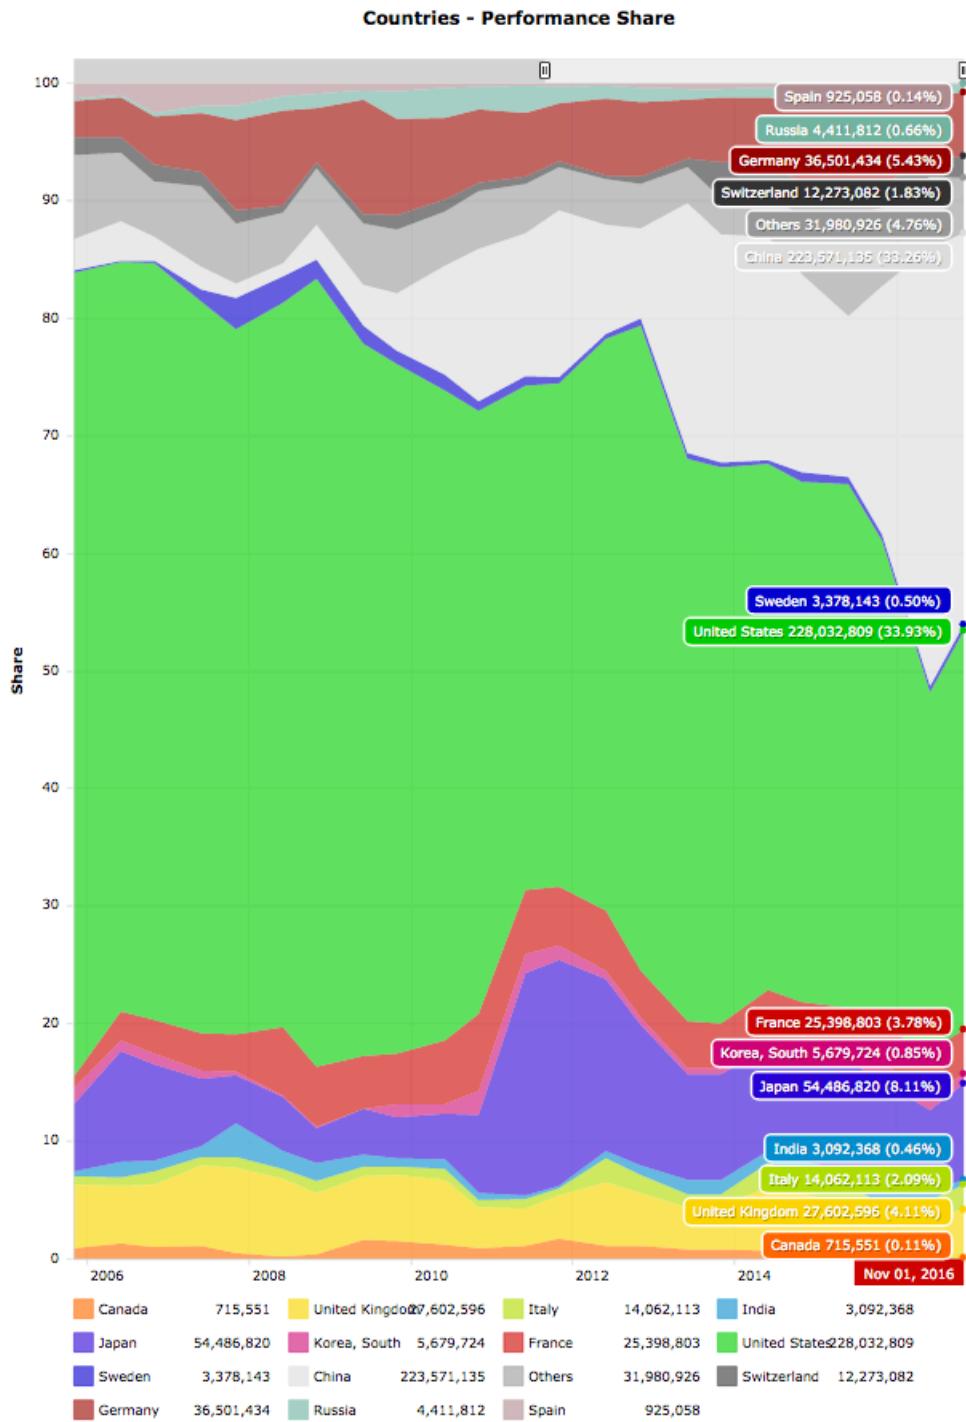
Rank	Site	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	National Supercomputing Center in Wuxi China	<b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway NRCPC	10,649,600	93,014.6	125,435.9	15,371
2	National Super Computer Center in Guangzhou China	<b>Tianhe-2 (MilkyWay-2)</b> - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P NUDT	3,120,000	33,862.7	54,902.4	17,808
3	DOE/SC/Oak Ridge National Laboratory United States	<b>Titan</b> - Cray XK7 , Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x Cray Inc.	560,640	17,590.0	27,112.5	8,209
4	DOE/NNSA/LLNL United States	<b>Sequoia</b> - BlueGene/Q, Power BQC 16C 1.60 GHz, Custom IBM	1,572,864	17,173.2	20,132.7	7,890
5	DOE/SC/LBNL/NERSC United States	<b>Cori</b> - Cray XC40, Intel Xeon Phi 7250 68C 1.4GHz, Aries interconnect Cray Inc.	622,336	14,014.7	27,880.7	3,939
6	Joint Center for Advanced High Performance Computing Japan	<b>Oakforest-PACS</b> - PRIMERGY CX1640 M1, Intel Xeon Phi 7250 68C 1.4GHz, Intel Omni-Path Fujitsu	556,104	13,554.6	24,913.5	2,719

## Projected Performance Development



*Lists*

● Sum    ▲ #1    ■ #500



Stanford University