



CME 213
SPRING 2017


Eric Darve

ANNOUNCEMENT: PROF. JINCHAO XU'S 1-UNIT CLASS

- **CME 335A Optimal Iterative Methods for Linear and Nonlinear Problems**
- **1-unit course meeting on 4/10, 4/12, 4/14, 4/17 and 4/19 from 3:30-5:20pm in GESB150.**
- **Instructor: Prof. Jinchao Xu, Verne M. Willaman
Professor of Mathematics, Penn State University**
- **Prof. Xu is best known for an algorithm that is now one of the two most fundamental multigrid approaches for solving large-scale PDEs – the Bramble-Pasciak-Xu preconditioner – and one of the most efficient methods for solving Maxwell's equations – the Hiptmair-Xu preconditioner.**

HOMWORK INSTRUCTIONS

- Download homework handout and skeleton code from canvas.
- Copy all the files to **corn**
- Submit using the provided script
- Turn in computer code + PDF with text answers.
- Deadline is: Wednesday April 12th, 11pm
- There is a 24 hour grace period.




The background of the slide features a large, light gray watermark of the Stanford University seal. The seal is circular and contains a redwood tree in the center, with the words "LELAND STANFORD JUNIOR UNIVERSITY" around the top and "1891" at the bottom. The text "EXAMPLE OF PARALLEL COMPUTATION" is centered over the seal.

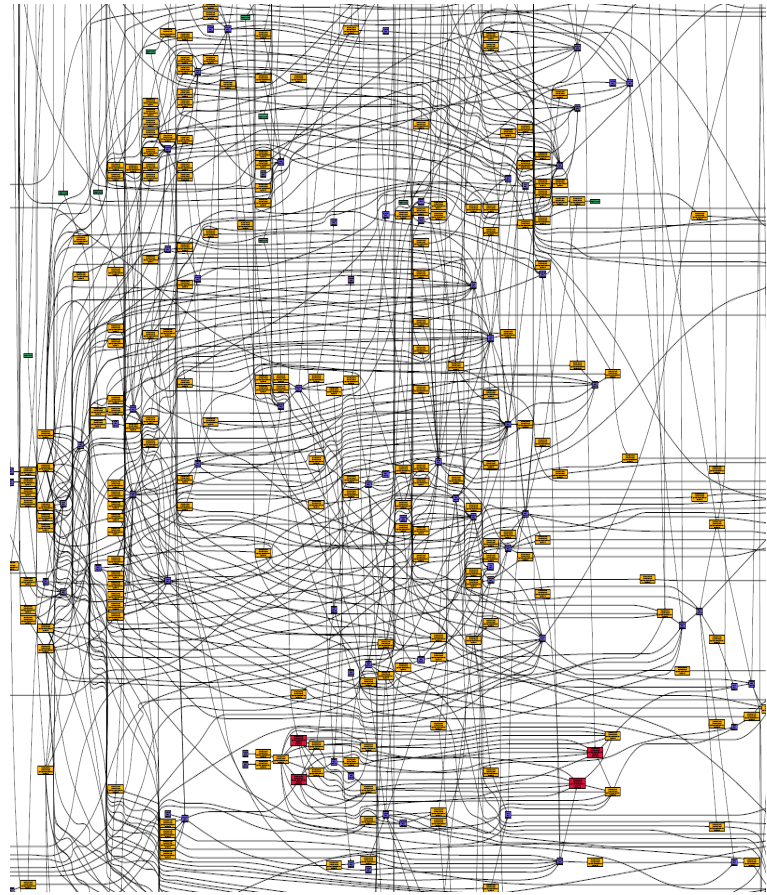
EXAMPLE OF PARALLEL COMPUTATION

WHY WE NEED TO WRITE PARALLEL PROGRAMS

- Most programs you have written so far are (probably) sequential.
- Unfortunately parallel programs often look very different...



Sequential program



Parallel program

- **An efficient parallel implementation of a serial program may not be obtained by simply parallelizing each step.**
- **Rather, the best parallelization may be obtained by stepping back and devising an entirely new algorithm.**

LET'S CALCULATE THE SUM OF N NUMBERS

```
sum = 0;
for (i = 0; i < n; i++) {
    x = ComputeNextValue(...);
    sum += x;
}
```


OUR FIRST PARALLEL PROGRAM

- Assume we have p cores that can compute and exchange data.
- Then we could accelerate the previous calculation by splitting the work among all these cores.


```
my_sum = 0;
my_first_i = ... ;
my_last_i = ... ;
for (my_i = my_first_i; my_i < my_last_i; my_i++) {
    my_x = ComputeNextValue(...);
    my_sum += my_x;
}
```

BUT IT'S NOT THAT SIMPLE

- Each core has computed a partial sum.
- All these partial sums need to be summed up together.
- The simplest approach is to have one “master” core do all the work:

In pseudo-code:

```
if (I am the master core) {  
    sum = my_x;  
    for each core other than myself {  
        receive value from core;  
        sum += value;  
    }  
} else {  
    send my_x to the master;  
}
```




THAT MAY NOT BE ENOUGH

- If we have many cores, this final sum may in fact take a lot of time.



- How would you design a better implementation?



AUTOMATIC PARALLELIZATION


- This simple example illustrates the fact that it is difficult for a compiler to parallelize a program.
- Instead the programmer must often re-write his code having in mind that multiple cores will be computing in parallel.
- The purpose of this class is to teach you the most common parallel languages used in science and engineering.

The background of the slide features a large, light gray watermark of the Stanford University seal. The seal is circular, with a diamond-patterned border. Inside the border, the text "LELAND STANFORD JUNIOR UNIVERSITY" is written in a circular path. Below this, the German phrase "DIE LUFT DER FREIHEIT WEHT" is visible. In the center of the seal is a detailed illustration of a redwood tree standing on a rocky outcrop. At the bottom of the seal, the year "1891" is inscribed.

SHARED MEMORY PROCESSOR

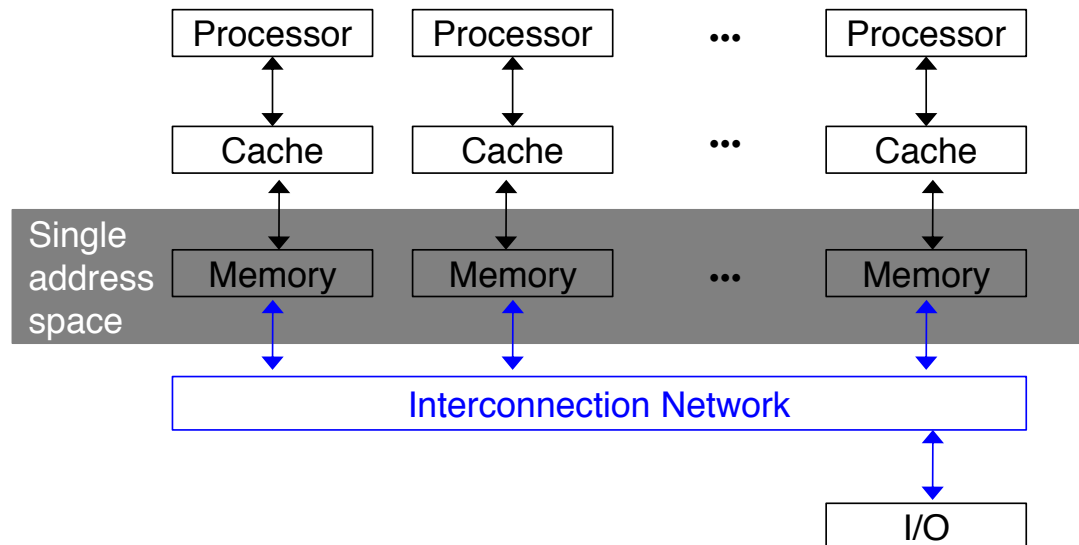
SCHEMATIC OF A MULTICORE PROCESSOR

- Model for shared memory machines
- Comprised of:
 - A number of processors or cores
 - A shared physical memory (global memory)
 - An interconnection network to connect the processors with the memory.

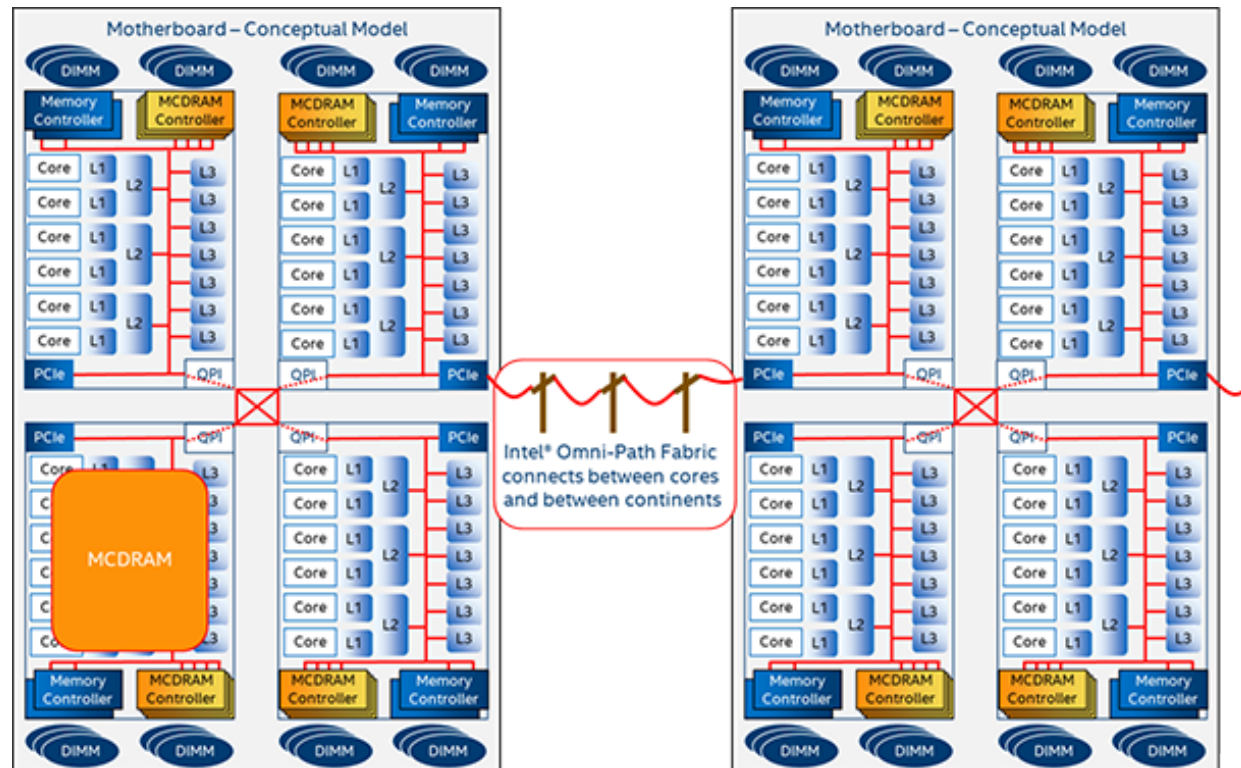


SHARED MEMORY NUMA

- In many cases, the program views the memory as a single addressable space, but in reality the memory is physically distributed.
- NUMA: non-uniform memory access.
- Faster access to memory, but special hardware required to move data between memory banks, e.g., Intel Omni-Path Fabric.



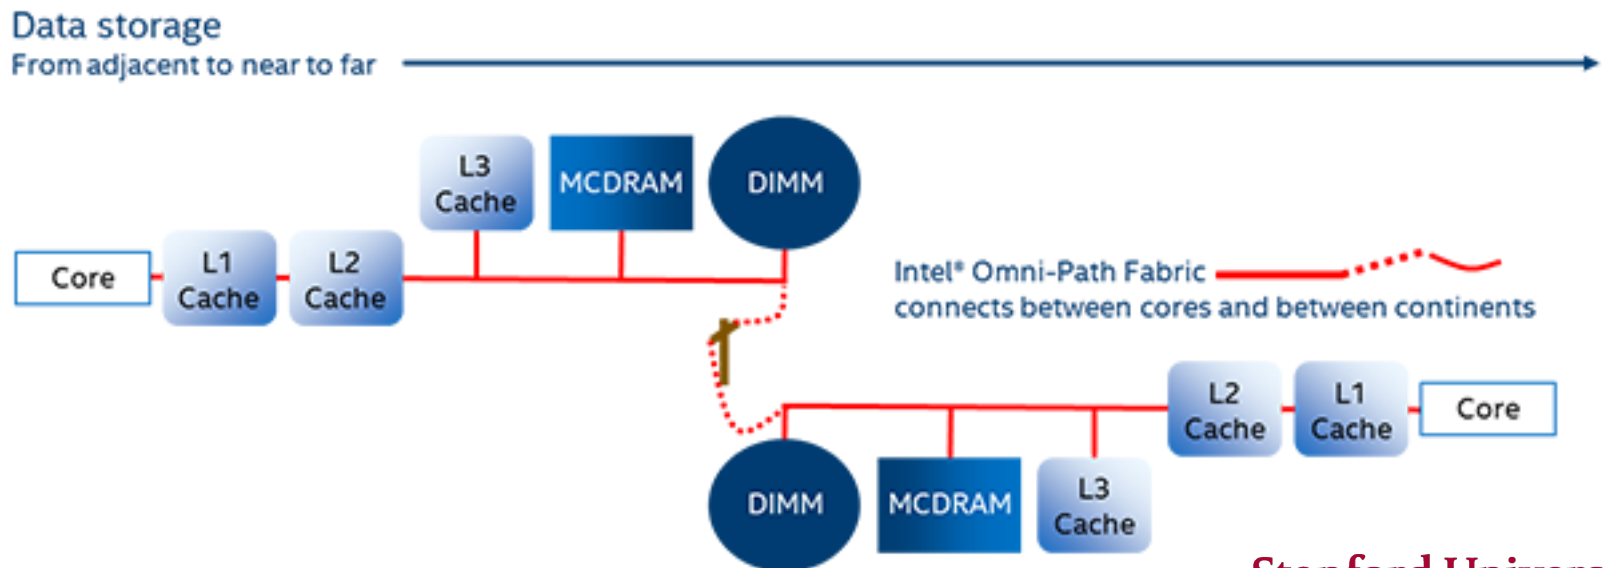
CASE STUDY: INTEL PROCESSOR



- MCDRAM: proprietary, high-bandwidth memory that physically sits atop Xeon Phi processors (Knights Landing). HBM: Xeon Phi Knights Hill.
- Omni-Path connects a core to memory sitting next to other cores.

MEMORY IS HIERARCHICAL

- In this class we will only briefly discuss performance for multicore processors.
- Things to keep in mind for performance:
 - memory is key to developing high-performance multicore applications
 - memory is hierarchical and complex.




SIZE, LATENCY, BANDWIDTH OF MEMORY SUBSYSTEM

Memory	Size	Latency	Bandwidth
L1 cache	32 KB	1 nanosecond	1 TB/second
L2 cache	256 KB	4 nanoseconds	1 TB/second Sometimes shared by two cores
L3 cache	8 MB or more	10x slower than L2	>400 GB/second
MCDRAM		2x slower than L3	400 GB/second
Main memory on DDR DIMMs	4 GB-1 TB	Similar to MCDRAM	100 GB/second
Main memory on Intel Omni-Path Fabric	Limited only by cost	Depends on distance	Depends on distance and hardware
I/O devices on memory bus	6 TB	100x-1000x slower than memory	25 GB/second
I/O devices on PCIe bus	Limited only by cost	From less than milliseconds to minutes	GB-TB/hour Depends on distance and hardware

PROCESSES AND THREADS


Definition:

- **Process:**
 - Program in execution.
 - **Comprises:** the executable program along with all information that is necessary for the execution of the program.
- **Thread:** an extension of the process model. Can be viewed as a “lightweight” process.
- In this model, each process may consist of multiple independent control flows that are called threads.
- A thread may be described as a “procedure” that runs independently from the main program.
- Imagine a program that contains a number of procedures. Then imagine these procedures being able to be scheduled to run simultaneously and/or independently by the operating system. This describes a **“multi-threaded” program**.



SHARED ADDRESS SPACE


- All the threads of one process share the address space of the process, i.e., they have a common address space.
- When a thread stores a value in the shared address space, another thread of the same process can access this value.



A large, light gray watermark of the Stanford University seal is centered in the background. The seal is circular with a diamond-patterned border. Inside the border, the text "LELAND STANFORD JUNIOR UNIVERSITY" is written in a circular path. Below this, the German phrase "DIE LUFT DER FREIHEIT WEHT" is written. In the center of the seal is a detailed illustration of a redwood tree standing on a rocky outcrop. At the bottom of the seal, the year "1891" is inscribed.

PTHREADS

THREADS ARE EVERYWHERE



- C threads: **Pthreads**
- C++ threads (11):
`std::thread`
- Java threads:
`Thread thread = new Thread();`
- Python threads:
`t = threading.Thread(target=worker)`
- Cilk:
`x = spawn fib (n-1);`
- Julia:
`r = remotecall(rand, 2, 2, 2)`
- **OpenMP**

PTHREADS

- This is the most “low-level” approach for programming in parallel.
- **Pthreads:** POSIX threads. This is a standard to implement threads on UNIX systems. It is based on the C programming language.
- Pthreads will serve as an introduction to the most important concepts in multicore programming.
- The other approach we will cover is OpenMP.
- OpenMP is the standard for multicore programming in scientific programs.
- Pthreads will help you understand OpenMP.



THE BASICS

- Include the header file:

```
include <pthread.h>
```

- Compile using:

```
gcc -o hello_pthread hello_pthread.c  
-lpthread
```

- See `hello_pthread.c`

THREAD CREATION

```
int pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*routine)(void*),  
    void *arg)
```

- **thread** thread identifier
- **routine** function that will be executed by the thread
- **arg** pointer to the argument value with which the thread function **routine()** will be executed
- **attr** use **NULL** for the time being

THREAD TERMINATION

A thread terminates when:

1. Thread reaches the end of its thread function, i.e., returns.
2. Thread calls

```
void pthread_exit(void *valuep)
```

Note:

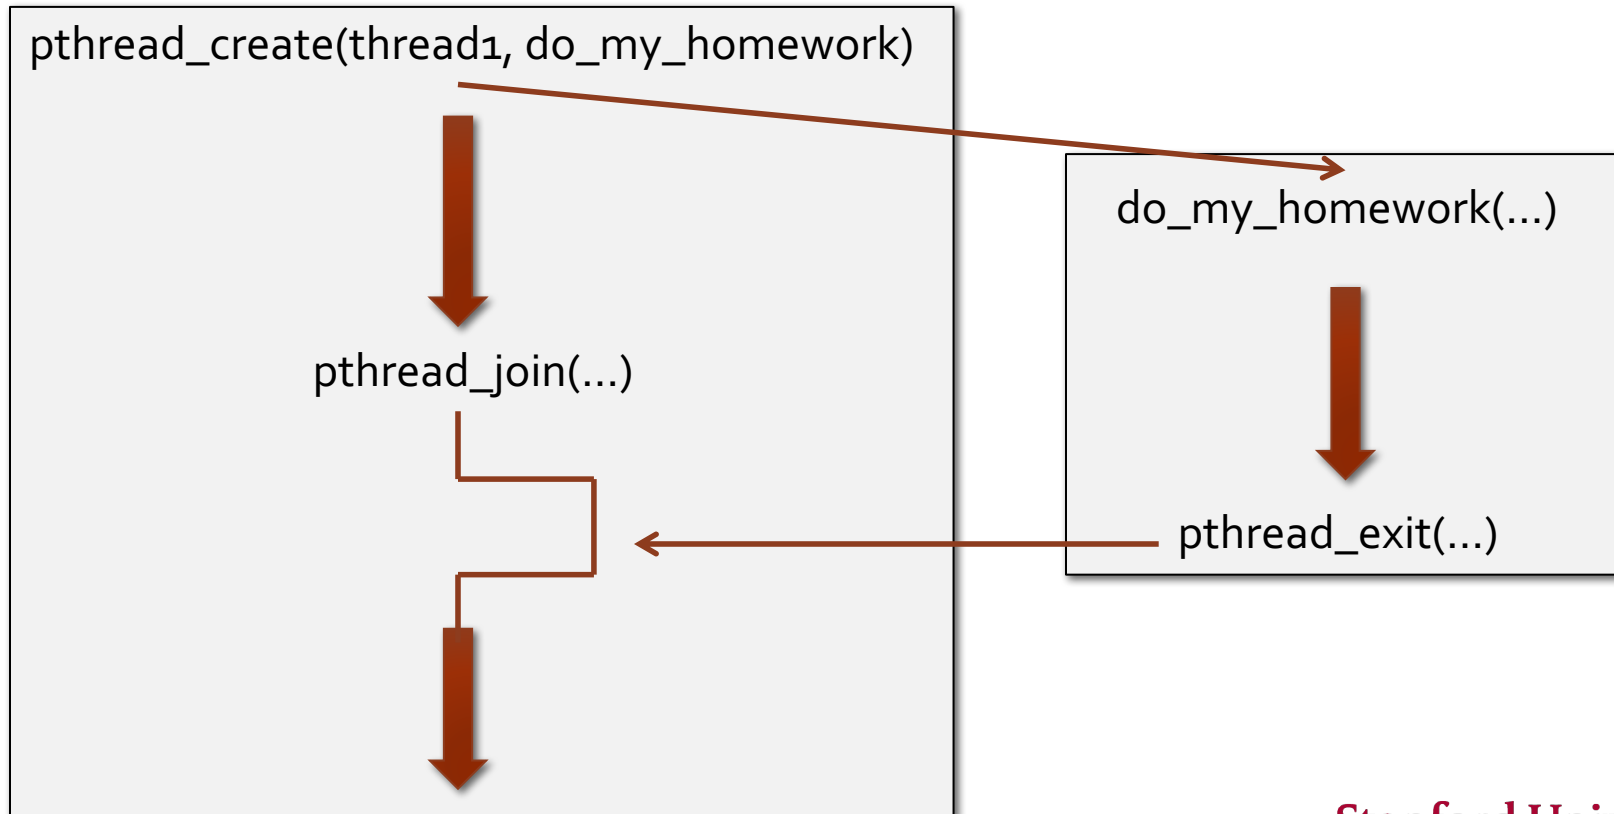
- Upon termination, a thread releases its runtime stack.
- Therefore, the pointer should point to: 1) a global variable, or 2) a dynamically allocated variable.

```
int pthread_join(pthread_t thread, void **valuep)
```

- Calling thread waits for `thread` to terminate.
- `pthread_join` is used to synchronize threads.
- `valuep` memory address where the return value of thread will be stored.

Thread 0

Thread 1



See

`hello_pthread_bug_1.c`

`hello_pthread_bug_2.c`



computer, shared memory is physically in the
location

True

False

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0%

`read_exit(valuep): valuep is a pointer pointing`

A global
variable

A local variable

A dynamically
allocated
variable

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0%

What is the role of pthread_join?

Terminate another thread

Check whether another thread has...

Create a new thread

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at Pollev.com/app

0%

What's the error in bug_1.c?

The type of result is wrong

The type of p_thread_result is wrong

result is a local variable

result no longer exists after the thread terminates

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at Pollev.com/app

0%

Line would you delete in order for bug_2.c to

Line

49

Line

51

Line

52

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0%

What happens if you delete line 49 in bug_2.c

The result is
correct

The result is
wrong

The result is
undetermined

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PolleEv.com/app

0%