

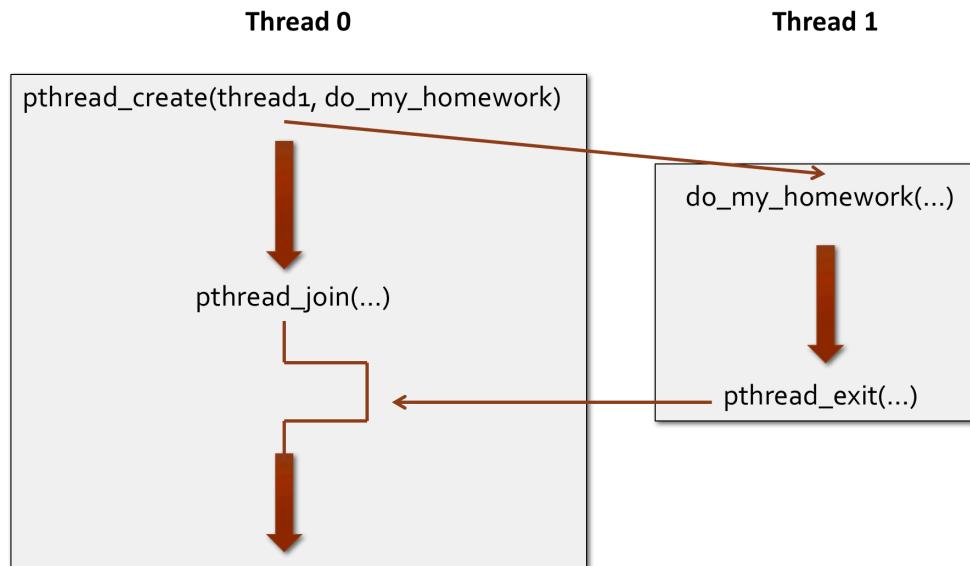
CME 213
SPRING 2017

Eric Darve

Stanford University

PTHREADS

- Low-level API
- Implemented in C
- `pthread_create`
- `pthread_exit`
- `pthread_join`



See

hello_pthread_bug_1.c

hello_pthread_bug_2.c



What's the error in bug_1.c?

The type of result is wrong

The type of p_thread_result is wrong

result is a local variable

result no longer exists after the thread terminates

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

Line would you delete in order for bug_2.c to

Line

49

Line

51

Line

52

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0%

What happens if you delete line 49 in bug_2.c

The result is
correct

The result is
wrong

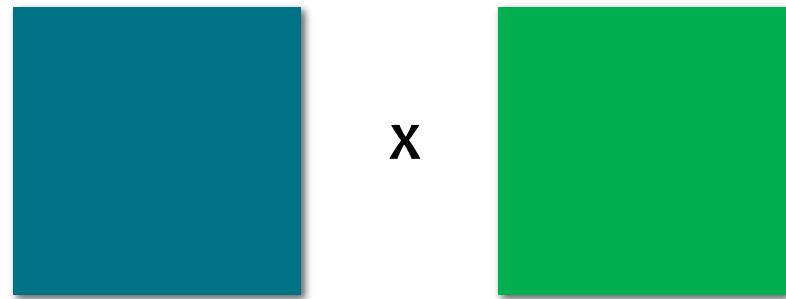
The result is
determined

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

EXAMPLE: MATRIX MATRIX PRODUCT

- $C = A * B$
where A and B are square matrices.
- See `matrix_prod.c`



Why are we casting work to void* in line 179?

↳ MatrixMult can be a general function

Because this is the type that pthread_create expects

Because this is required by the

future of MatrixMult

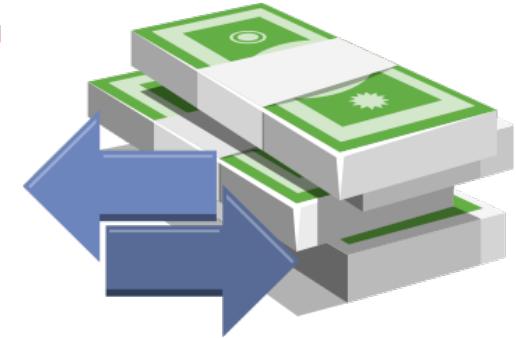
Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

THREAD COORDINATION



THE RISKS OF MULTI-THREADED PROGRAMMING



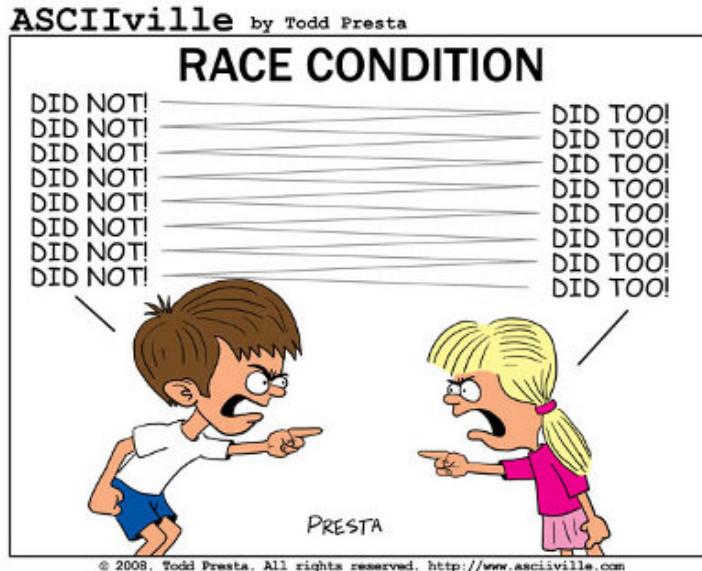
- Let us assume that a well-known bank company has asked you to implement a multi-threaded code to perform bank transactions.
- You start with the modest goal of allowing deposits.
- Clients deposit money and the amount gets credited to their accounts.
- As a result of having multiple threads running concurrently the following can happen:

A PARALLEL BANK DEPOSIT

Thread 0	Thread 1	Balance
Client requests a deposit	Client requests a deposit	\$1000
Check current balance = \$1000	Check current balance = \$1000	
Ask for deposit amount = \$100	Ask for deposit amount = \$300	
	Compute new balance = \$1300	
Compute new balance = \$1100	Write new balance to account	\$1300
Write new balance to account		\$1100

RACE CONDITION

- Although the correct balance should be \$1,400, it is \$1,100.
- The problem is that many operations “take time” and can be interrupted by other threads attempting to modify the same data.
- This is called a race condition: the final result depends on the precise order in which the instructions are executed.



RACE CONDITION

- This issue is addressed using mutexes (mutex): mutual exclusion.
- They ensure that certain common pieces of data are accessed and modified by a single thread.
- This problem typically occurs when you have a sequence like:
READ/WRITE, or
WRITE/READ
performed by different threads.

Thread 0 wants
to add new to-do
item.



Thread 0 closes lock.
No other thread can
open the lock.



Thread 0 is done
with the to-do list.
It opens the lock.

Thread 1 wants to
open the lock.
It has to wait.



Thread 1 can close the
lock and access the to-
do list.



MUTEX

- A mutex can only be in two states: locked or unlocked.
- Once a thread locks a mutex:
 - Other threads attempting to lock the same mutex are blocked
 - Only the thread that initially locked the mutex has the ability to unlock it.
- This allows to protect regions of code.



TYPICAL USAGE

Mutex use:

- Create and initialize a mutex variable.
- Threads attempt to lock the mutex.
- Only one succeeds, and that thread owns the mutex.
- The owner thread performs some set of actions.
- The owner unlocks the mutex.
- Another thread acquires the mutex and repeats the process.



PIZZA RESTAURANT



Pizza delivery team



- Checks the addresses for customers
- Deliver pizza

Go back;
check if there
are orders left.



PIZZA RESTAURANT CODE

See

mutex_demo.c

SUMMARY OF KEY FUNCTIONS

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

```
int pthread_mutex_init(pthread_mutex_t *mutex,  
                      const pthread_mutexattr_t *attr)
```

Initialization of mutex; choose NULL for attr.

```
int pthread_mutex_destroy(pthread_mutex_t *mutex)
```

Destruction of mutex.

```
int pthread_mutex_lock(pthread_mutex_t *mutex)
```

Locks a mutex; blocks if another thread has locked this mutex and owns it.

```
int pthread_mutex_unlock(pthread_mutex_t *mutex)
```

Unlocks mutex; after unlocking, other threads get a chance to lock the mutex.

Should I move line 76 (`mutex_lock`) to before line
"global_task_list = global_task_list->next")

Yes because the thread
is not modifying
global_task_list until
line 86

No because of line 79

It depends on what
threads are doing

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

Yes this is safer

No this will make
the code go
slower

It makes no
difference

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

DEADLOCK

Another strange case (but common!) of parallel computing.

Let's assume:

Thread 0 Locks *mutex0*

Thread 1 Locks *mutex1*

Thread 0 Does some work

Thread 1 Does some work

Thread 0 Locks *mutex1*

Thread 1 Locks *mutex0*

The code will
never reach
this point



Thread 0 Work requiring lock on both mutexes

Thread 1 Work requiring lock on both mutexes

Thread 0 Unlocks all mutexes

Thread 1 Unlocks all mutexes



Solution: threads always lock mutexes in the same order.

CONDITION VARIABLES

ARE WE THERE YET?

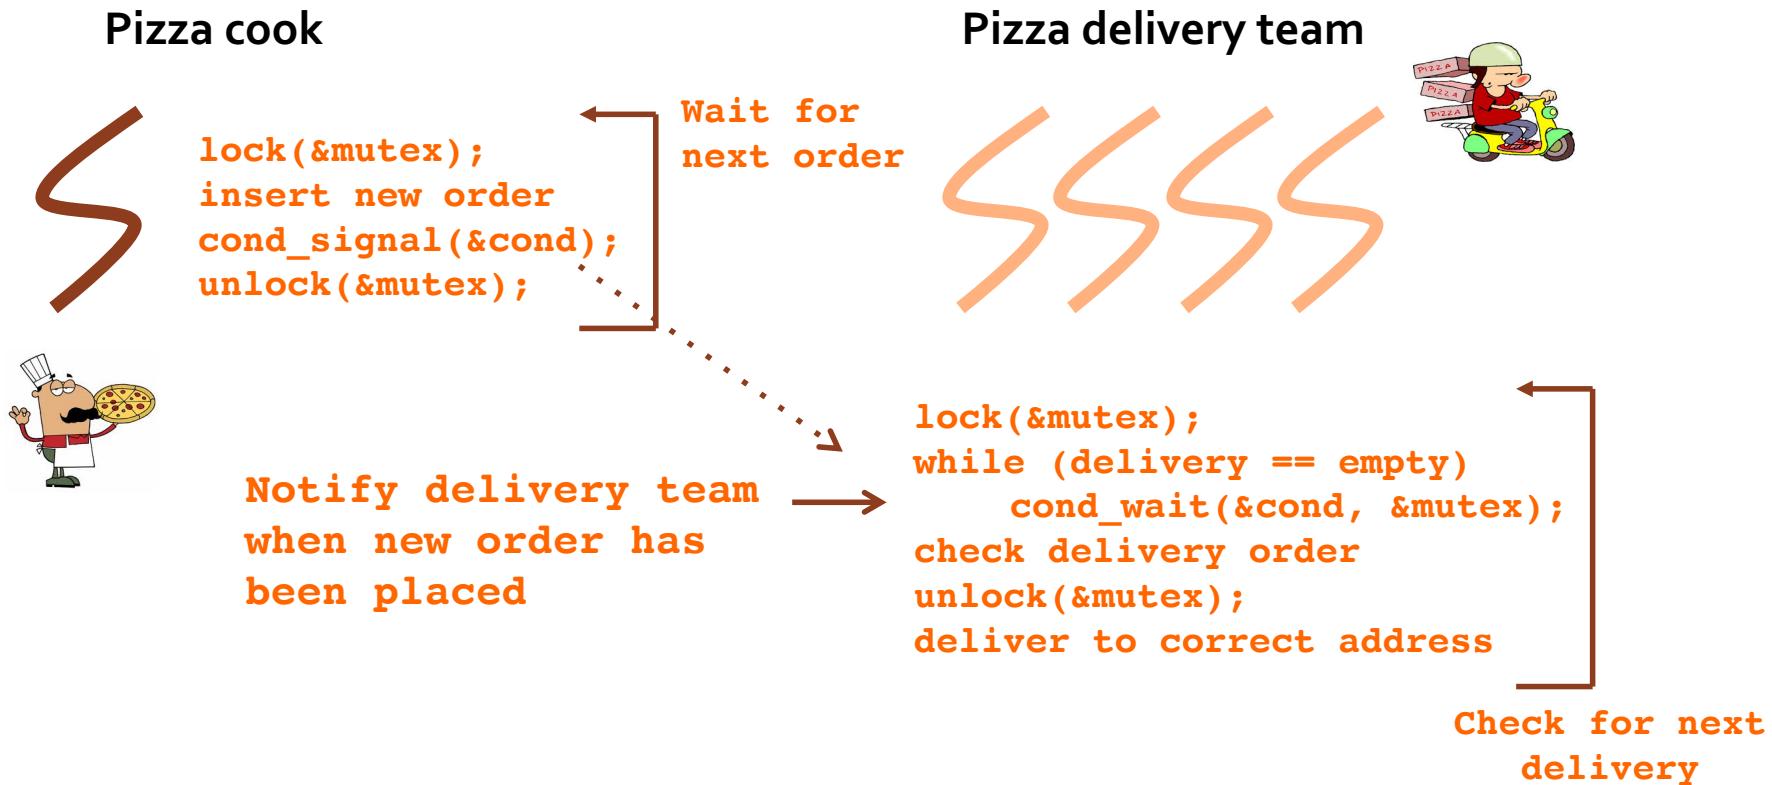


- There is a common situation in computing: you have a pool of threads that are ready to perform tasks, but there is nothing to do yet. So they have to wait until some work becomes available.
- Instead of constantly checking the to-do list, they can simply wait and be awoken when a certain condition is met.



WAITING ON A CONDITION

- Application: pizza delivery boys waiting on pizza cook.
- This is the producer-consumer model.



CONDITION VARIABLE EXAMPLE

cond_var.c

Is the mutex unlocked when the thread enters `pthread_cond_wait`?

Yes

No

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0%

opens to the mutex after a thread receives a signal?

It is
unlocked

It is
locked

Nothing

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

CONDITION VARIABLE

- Requires a condition variable and a mutex.
- The mutex is used to protect the access to the condition variable.

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex)
```

Recommended scenario

```
pthread_mutex_lock(&mutex);  
while (!condition_ready())  
    pthread_cond_wait(&cond, &mutex);  
access_modify_shared_data();  
pthread_mutex_unlock(&mutex);
```

When `pthread_cond_wait` is called:

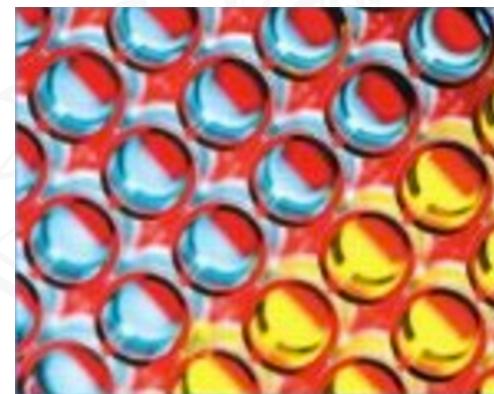
- The thread releases the mutex `mutex`.
- The thread waits until a signal is sent.
- Upon receiving a signal, the thread locks `mutex` and proceeds.

CONDITION SIGNAL

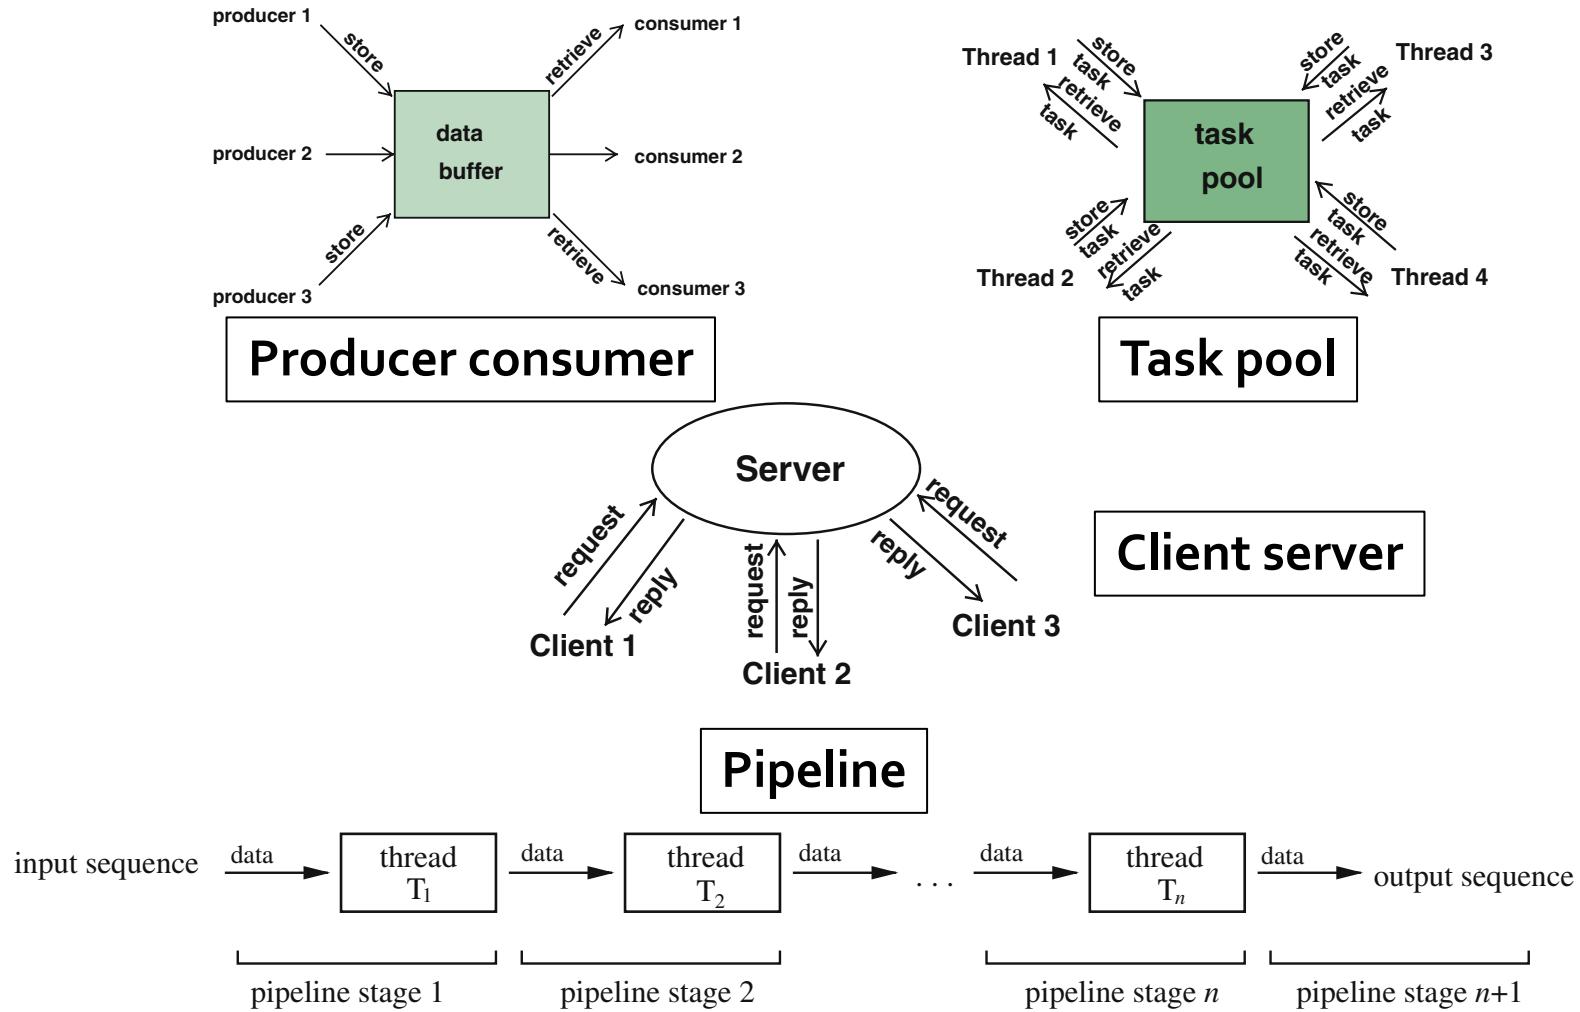
```
int pthread_cond_signal(pthread_cond_t *cond)
```

- Wakes up a single thread waiting on the variable **cond**.
- Typically a mutex is used around **pthread_cond_signal** to protect the evaluation of the condition.

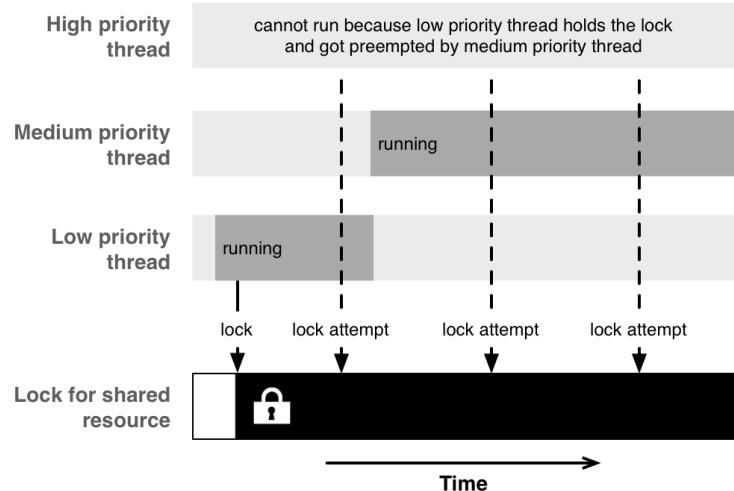
PARALLEL PATTERNS AND OTHER FEATURES OF PTHREADS



EXAMPLE OF USEFUL PARALLEL PROGRAMMING PATTERNS



OTHER FEATURES



- **Thread scheduling**
 - Implementations will differ on how threads are scheduled to run. In most cases, the default mechanism is adequate.
 - The Pthreads API provides routines to explicitly set thread scheduling policies and priorities which may override the default mechanisms.
- **Thread-specific data: keys**
 - To preserve stack data between calls, you can pass it as an argument from one routine to the next, or else store the data in a global variable associated with a thread.
 - Pthreads provides another, possibly more convenient and versatile, way of accomplishing this, through keys.
- **Priority inversion problems; priority ceiling and priority inheritance**
- **Thread cancellation**
- **Barriers; not always available**
- **Thread safety**