

A large, light gray watermark of the Stanford University seal is centered in the background. The seal is circular with a diamond-patterned border. Inside the border, the text "LELAND STANFORD JUNIOR UNIVERSITY" is written in a circular path. Below this, the German phrase "DIE LUFT DER FREIHEIT WEHT" is written. In the center of the seal is a redwood tree standing on a hill. At the bottom of the seal, the year "1891" is inscribed.

CME 213

SPRING 2017

Eric Darve

Review

- “Secret” behind GPU performance: simple cores but a large number of them; even more threads can exist live on the hardware (10k–20k threads live).
- Important performance bottleneck: data traffic between memory and register files
- Programming on a GPU:
 - Warp = 32 threads
 - Block = group of threads running on one SM
 - SM runs an integer number of thread blocks
 - Grid = set of blocks representing the entire data set

Code overview

```
int* d_output;
cudaMalloc(&d_output, sizeof(int) * N);
kernel<<<1, N>>>(d_output);
vector<int> h_output(N);
cudaMemcpy(&h_output[0], d_output, sizeof(int) * N,
           cudaMemcpyDeviceToHost);
for(int i = 0; i < N; ++i) {
    printf("Entry %3d, written by thread %2d\n",
           h_output[i], i);
}
cudaFree(d_output);
```

Function to execute

Input variables

`kernel<<<1, N>>>(d_output);`

Number of threads to use

Device kernels

```
__device__ __host__  
int f(int i) {  
    return i*i;  
}
```

```
__global__  
void kernel(int* out) {  
    out[threadIdx.x] = f(threadIdx.x);  
}
```

global host device

What are these mysterious keywords?

`__global__` kernel will be

- Executed on the device
- Callable from the host

`__host__` kernel will be

- Executed on the host
- Callable from the host



Compiled for host

`__device__` kernel will be

- Executed on the device
- Callable from the device only



Compiled for device

```
./firstProgram 1  
./firstProgram 32  
./firstProgram 1024  
./firstProgram 1025
```

- Fails at 1025!
- Blocks and grid required for “real” runs.



Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app



You can use `__global__` and `__host__` together

Yes, the
function gets
compiled...

No, these
options are

incompatible

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0

Total Results: 0

Thread hierarchy

Two levels of hierarchy:

- **Thread block:** a block is loaded entirely on an SM. This means that a block of threads cannot require more resources than available, for example to store register variables.
- **Grid of blocks:** for large problems, we can use multiple blocks. A very large number of blocks can be allocated so that we can execute a kernel over a very large data set.

Problem decomposition

Decompose your problem into:

- Small blocks so that data can fit in SM.
 - Typically the number of threads per block is around 256–512.
- Create a grid of blocks so that you can process the entire data.
- Hardware limitations for GPUs on Certainty:

Maximum threads per block:	1024
Maximum dimension 0 of block:	1024
Maximum dimension 1 of block:	1024
Maximum dimension 2 of block:	64
Maximum dimension 0 of grid:	65535
Maximum dimension 1 of grid:	65535
Maximum dimension 2 of grid:	65535

Multithreaded CUDA Program

Block 0

Block 1

Block 2

Block 3

Block 4

Block 5

Block 6

Block 7

GPU with 2 SMs

SM 0

SM 1

Block 0

Block 1

Block 2

Block 3

Block 4

Block 5

Block 6

Block 7

GPU with 4 SMs

SM 0

SM 1

SM 2

SM 3

Block 0

Block 1

Block 2

Block 3

Block 4

Block 5

Block 6

Block 7

Block execution

The device will start by loading data on each SM to execute the blocks.

- At least one block must fit on an SM, e.g., there should be enough memory for register variables.
- The SM will load up as many blocks as possible until it runs out of resources.
- Then it will start executing the blocks, that is all threads will execute the kernel.
- Once all threads in a block are done with the kernel, another block gets loaded in the SM.
- And so on until the device runs out of blocks to execute.

Dimensions

Blocks and grids can be 1D, 2D or 3D.
Their dimensions is declared using:

```
dim3 threadsPerBlock(Nx);  
dim3 numBlocks(Mx);
```

```
dim3 threadsPerBlock(Nx, Ny);  
dim3 numBlocks(Mx, My);
```

```
dim3 threadsPerBlock(Nx, Ny, Nz);  
dim3 numBlocks(Mx, My, Mz);
```

Thread index example:

```
int col = blockIdx.x * blockDim.x + threadIdx.x;  
int row = blockIdx.y * blockDim.y + threadIdx.y;
```

Block ID



Block size

Thread ID in block

addMatrices.cu

```
__global__
void Add(int n, int* a, int* b, int* c) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;

    if(i < n && j < n) {
        c[n*i + j] = a[n*i + j] + b[n*i + j];
    }
}

/* Dimension along y = n_thread; dimension along x and z = 1 */
dim3 threads_per_block(1, n_thread);

int blocks_per_grid_x = n;
int blocks_per_grid_y = (n + n_thread - 1) / n_thread;
dim3 num_blocks(blocks_per_grid_x, blocks_per_grid_y);

/* Run calculation on GPU */
Initialize<<<num_blocks, threads_per_block>>>(n, d_a, d_b);
Add<<<num_blocks, threads_per_block>>>(n, d_a, d_b, d_c);
```

Should I use a grid of blocks if I have a simple line
 $x[i]$ to compute with?

Yes

No

It depends

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0

Total Results: 0

possible to fill up all the SMs with a single huge

Yes, but it's
not
recommended

No, it won't
happen

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0

Total Results: 0

Compiler options and other tricks

- g** debug on the host
- G** debug on the device (see documentation, CUDA-gdb, Nsight Eclipse Edition)
- pg** profiling info for use with gprof
- xcompiler** (options for underlying gcc compiler)

More on profiling: **nvprof**

Profiling

- The visual profiler (**nvvp**) is the simplest way to profile your code. You get a graphical interface that helps navigate the information.
- You can also use a command line tool: **nvprof**.
- This can also be used to collect information on a cluster to be visualized using **nvvp** on your local computer.

- Run

```
nvprof ./addMatrices -n 4000
```

==9229== NVPROF is profiling process 9229, command: ./addMatrices -n 4000

==9229== Dimensions of matrix: 4000 x 4000

Dimension of thread block along y: 512

Dimension of grid: 4000 x 8

All tests have passed; calculation is correct.

Profiling application: ./addMatrices -n 4000

==9229== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
82.99%	17.713ms	1	17.713ms	17.713ms	17.713ms	[CUDA memcpy DtoH]
9.09%	1.9394ms	1	1.9394ms	1.9394ms	1.9394ms	Add(int, int*, int*, int*)
7.92%	1.6907ms	1	1.6907ms	1.6907ms	1.6907ms	Initialize(int, int*, int*)

==9229== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
80.16%	104.98ms	3	34.993ms	281.33us	104.37ms	cudaMalloc
13.77%	18.035ms	1	18.035ms	18.035ms	18.035ms	cudaMemcpy
2.79%	3.6503ms	2	1.8251ms	1.7031ms	1.9472ms	cudaDeviceSynchronize
2.67%	3.4989ms	332	10.538us	157ns	395.13us	cuDeviceGetAttribute
0.31%	403.71us	4	100.93us	100.75us	101.32us	cuDeviceTotalMem
0.25%	331.28us	4	82.819us	79.806us	88.013us	cuDeviceGetName
0.03%	40.453us	2	20.226us	12.109us	28.344us	cudaLaunch
0.00%	5.8130us	7	830ns	196ns	4.1910us	cudaSetupArgument
0.00%	3.1480us	2	1.5740us	329ns	2.8190us	cudaGetLastError
0.00%	2.4130us	2	1.2060us	274ns	2.1390us	cuDeviceGetCount
0.00%	2.2930us	2	1.1460us	361ns	1.9320us	cudaConfigureCall
0.00%	1.6810us	8	210ns	165ns	297ns	cuDeviceGet

cuda-memcheck

- Several tools available to verify your code.
- 4 tools accessible with `cuda-memcheck` command:
 - `memcheck`: Memory access checking
 - `initcheck`: Global memory initialization checking
 - `racecheck`: Shared memory hazard checking
 - `synccheck`: Synchronization checking;
`__syncthreads()`, not on SM 2.x
- Usage example:
`cuda-memcheck ./memcheck_demo`
`cuda-memcheck --leak-check full ./memcheck_demo`

GPU Optimization!



Top factors to consider

- **Memory: data movement!**
- **Occupancy and concurrency:**
 - “hide latency through concurrency”**
- **Control flow: branching**

Memory

- The number one bottleneck in scientific applications are not flops!
- Flops are plentiful and abundant.
- The problem is data starvation.
- Computing units are idle because they are waiting on data coming from / going to memory.

Memory access

- Similar to a regular processor but with some important differences.
- Caches are used to optimize memory accesses: L1 and L2 caches.
- Cache behavior is complicated and depends on the compute capability of the card (that is the type of GPU you are using).
- Generations (compute capability): 2.x, 3.x, 5.x, and 6.x (no 4.x).
- For this class, we have access to 2.0 hardware.
- We will focus on this for our discussion.

Cache

- There are two main levels of cache: L1 and L2.
- L2 is shared across SMs.
- Each L1 is assigned to only one SM.
- When reading data from global memory, the device can only read **entire cache lines**.
- There is no such thing as reading a single float from memory.
- To read a single float, you read the entire cache line and “throw out” the data you don’t need.

Cache lines

- **L1: 128-byte lines = 32 floats.**
- L2: 32-byte segments = 8 floats.
- Commit these fundamental numbers to memory now!

How is data read from memory then?

- Remember how the hardware works.
- All threads in a warp execute the same instruction at the same time.
- Take 32 threads in a warp.
- Each thread is going to request a memory location.
- Hardware groups these requests into a number of **cache line requests**.
- Then data is read from memory.
- At the end of the day, what matters is not how much data threads are requesting from memory, it's **how many memory transactions (cache lines)** are performed!

Example

```
__global__ void offsetCopy(float* odata, float* idata, int offset) {  
    int xid = blockIdx.x * blockDim.x + threadIdx.x + offset;  
    odata[xid] = idata[xid];  
}
```


offset = 0

This is the best case.

All memory accesses are coalesced into a single memory transaction.

One memory transaction


Aligned accesses (sequential/non-sequential)



	2.0 and later	
	Uncached	Cached
	1x 32B at 128 1x 32B at 160 1x 32B at 192 1x 32B at 224	1x 128B at 128

Misaligned by 1

Mis-aligned accesses (sequential/non-sequential)




Compute capability:	2.0 and later	
	Uncached	Cached
Memory transactions:	1x 32B at 128 1x 32B at 160 1x 32B at 192 1x 32B at 224 1x 32B at 256	1x 128B at 128 1x 128B at 256

```
__global__ void offsetCopy(float* odata, float* idata, int offset) {  
    int xid = blockIdx.x * blockDim.x + threadIdx.x + offset;  
    odata[xid] = idata[xid];  
}
```

offset = 1

Strided access

```
__global__ void stridedCopy(float* odata, float* idata, int stride) {  
    int xid = blockIdx.x * blockDim.x + threadIdx.x + offset;  
    odata[xid] = idata[stride*xid];  
}
```



- Stride is 2.
- Effective bandwidth of application reduced by 2.

Efficiency of reads

- Perfectly coalesced access: one memory transaction.
 - $\text{Bandwidth} = \text{Peak Bandwidth}$
- Misaligned: two memory transactions.
 - $\text{Bandwidth} = \text{Peak Bandwidth} / 2$
- Strided access
 - $\text{Bandwidth} = \text{Peak Bandwidth} / s$
for a stride of s
- Random access
 - $\text{Bandwidth} = \text{Peak Bandwidth} / 32$

```
__global__ void randomCopy(float* odata, float* idata, int* addr) {  
    int xid = blockIdx.x * blockDim.x + threadIdx.x + offset;  
    odata[xid] = idata[addr[xid]];  
}
```

Data type size

- The previous discussion covered the most common case where threads read a **4-byte float or int**.
- **char**: 1 byte. **float**: 4 bytes. **double**: 8 bytes.
- CUDA has many built-in data types that are shorter/longer:
 - char1-4, uchar1-4**
 - short1-4, ushort1-4**
 - int1-4, uint1-4**
 - long1-4, ulong1-4**
 - longlong1-2, ulonglong1-2**
 - float1-4**
 - double1-2**
- Example: **char1, int2, long4, double2**

size of a single memory transaction is 128 bytes

Only if the memory access is coalesced

No, this is only true when there is no offset in the access

Yes, if the type is float, but not double

Yes, when the bandwidth is close to peak

Yes

Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app

0

Total Results: 0

How to calculate the warp ID


- How do you know if 2 threads are in the same warp?
- Answer: look at the **thread ID**. Divide by 32: you get the warp ID.
- Thread ID is computed from the 1D, 2d or 3D thread index using:

$$x + yD_x + zD_xD_y$$

```
int tID = threadIdx.x
        + threadIdx.y * blockDim.x
        + threadIdx.z * blockDim.x * blockDim.y;
int warpID = tID/32;
```

Matrix transpose

- Let's put all these concepts into play through a specific example: a matrix transpose.
- It's all about bandwidth!
- Even for such a simple calculation, there are many optimizations.




simpleTranspose

```
template<typename T>
__global__
void simpleTranspose(T* array_in, T* array_out, int n_rows, int n_cols) {
    const int tid = threadIdx.x + blockDim.x * blockIdx.x;

    int col = tid % n_cols;
    int row = tid / n_cols;

    if(col < n_cols && row < n_rows) {
        array_out[col * n_rows + row] = array_in[row * n_cols + col];
    }
}
```

Memory access pattern



Coalesced reads



Strided writes

