

A large, light gray watermark of the Stanford University seal is centered in the background. The seal is circular with a diamond-patterned border. Inside the border, the words "LELAND STANFORD JUNIOR UNIVERSITY" are written in a circular path. Below this, the German phrase "DIE LUFT DER FREIHEIT WEHT" is written. In the center of the seal is a redwood tree standing on a hill. At the bottom of the seal, the year "1891" is inscribed.

CME 213

SPRING 2017

Eric Darve

The background of the slide features a large, light gray watermark of the Stanford University seal. The seal is circular and contains a redwood tree in the center, with the text "LELAND STANFORD JUNIOR UNIVERSITY" around the top and "DIE LUFT DER FREIHEIT WEHT" around the bottom. The year "1891" is at the very bottom. There are also stars around the inner circle.

LINEAR ALGEBRA

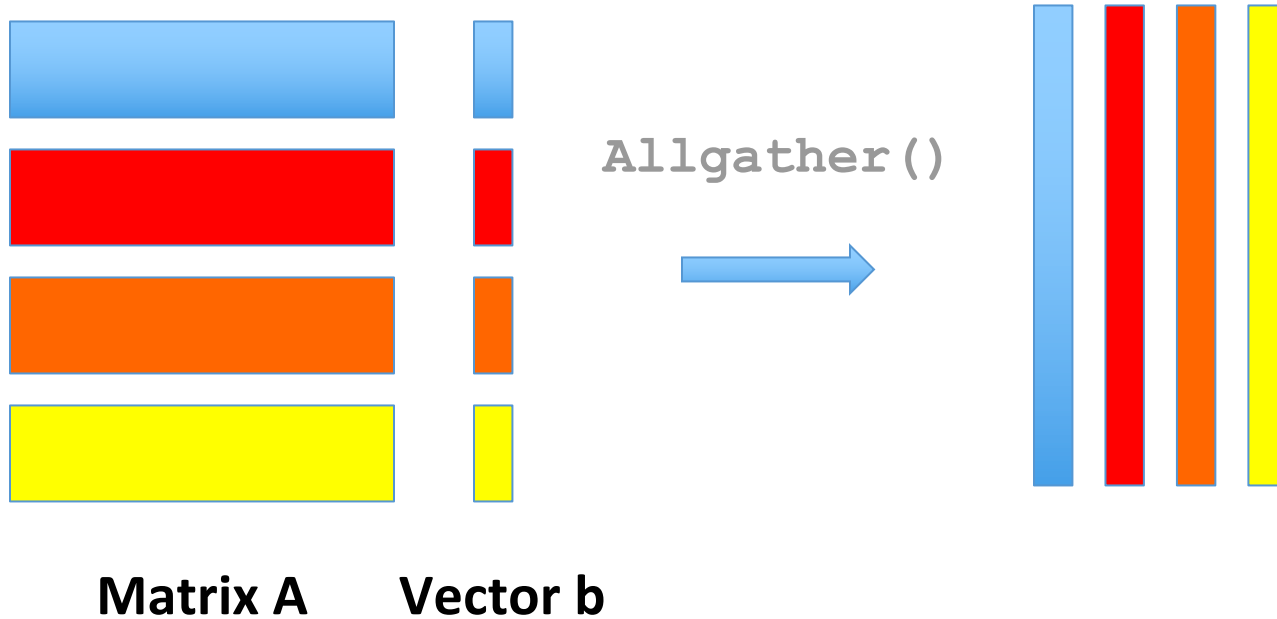
MATRIX-VECTOR PRODUCTS

Application example: matrix-vector product

- We are going to use that example to illustrate additional MPI functionalities.
- This will lead us to process groups and topologies.
- First, we go over two implementations that use the functionalities we have already covered.
- Two simple approaches:
 - Row partitioning of the matrix, or
 - Column partitioning

Row partitioning

This is the most natural.



Step 1: replicate b on each process: `MPI_Allgather()`

Step 2: perform product

See MPI code: `matvecrow/`


```
/* Gather entire vector b on each processor using Allgather */
MPI_Allgather(&bloc[0], nlocal, MPI_FLOAT, &b[0], nlocal, MPI_FLOAT,
| | | | | | | MPI_COMM_WORLD);
// sending nlocal and receiving nlocal from any other process

/* Perform the matrix-vector multiplication involving the
| locally stored submatrix. */
vector<float> x(nlocal);

for(int i=0; i<nlocal; i++) {
|   x[i] = 0.0;


|   for(int j=0; j<n; j++) {
| |   x[i] += a[i*n+j]*b[j];
| | }
| }
}
```

Column partitioning



Step 1: calculate partial products with each process


Column partitioning (cont'd)



- Step 2: reduce all partial results: `MPI_Reduce()`
- Step 3: send sub-blocks to all processes: `MPI_Scatter()`
- Steps are very similar to row partitioning.



Row and column partitioning schemes are dual of each other



Yes **A**

No **B**



Start the presentation to activate live content




If you see this message in presentation mode, install the add-in or get help at PolleEv.com/app

Total Results: 0


A better partitioning

- If the number of processes becomes large compared to the matrix size, we need a 2D partitioning:




- Each colored square can be assigned to a process.
- This allows using more processes.
- In addition, a theoretical analysis (more on this later) shows that this scheme runs faster.


Outline of algorithm: step 1



First column
contains b



Send b to the
diagonal
processes




Send b down each
column.

This is a broadcast operation.

Step 2 and 3

- Step 2: perform matrix-vector product locally
- Step 3: reduce across columns and store result in column 0.




Reduction across columns

Communication cost (in a nutshell)

Why is 2D partitioning better?




Reduction: $n/2$



Larger blocks

Reduction: $2n$



Narrow columns

Difficulties with 2D partitioning

- This type of decomposition brings some difficulties.
- We used two collective operations:
 - A broadcast inside a column.
 - A reduction inside a row.
- To do this in MPI, we need two concepts:
 - **Communicators or process groups.** This defines a subset of all the processes. For each subset, collective operations are allowed, e.g., broadcast for the group of processes inside a column.
 - **Process topologies.** For matrices, there is a natural 2D topology with (i,j) block indexing. MPI supports such grids (any dimension). Using MPI grids (called “Cartesian topologies”) simplifies many MPI commands.

The background of the slide features a large, light gray watermark of the Stanford University seal. The seal is circular and contains a redwood tree in the center, with the words "LELAND STANFORD JUNIOR UNIVERSITY" around the top and "1891" at the bottom. The text "PROCESS GROUPS AND COMMUNICATORS" is overlaid on the seal.

PROCESS GROUPS AND COMMUNICATORS

Process groups

- Groups are needed for many reasons.
- **Enables collective communication operations across a subset of processes.**
- Allows to easily assign independent tasks to different groups of processes.
- Provides a good mechanism to integrate a parallel library into an MPI code.

Groups and communicators

- **A group** is an ordered set of processes.
- Each process in a group is associated with a unique integer rank. Rank values start at zero and go to N-1, where N is the number of processes in the group.
- A group is always associated with a communicator object.
- **A communicator** encompasses a group of processes that may communicate with each other. All MPI messages must specify a communicator.
- For example, the handle for the communicator that comprises all tasks is `MPI_COMM_WORLD`.
- From the programmer's perspective, a group and a communicator are almost the same. The group routines are primarily used to specify which processes should be used to construct a communicator.
- Processes may be in more than one group/communicator. They have a unique specific rank within each group/communicator.

Main functions

MPI provides over 40 routines related to groups, communicators, and virtual topologies!

```
int MPI_Comm_group(MPI_Comm comm, MPI_Group *group)
```

Return group associated with communicator, e.g., MPI_COMM_WORLD

```
int MPI_Group_incl(MPI_Group group, int p, int *ranks,  
    MPI_Group *new_group)
```

ranks integer array with p entries.

Creates a new group `new_group` with p processes, which have ranks from 0 to p-1. Process i is the process that has rank `ranks[i]` in group.

```
int MPI_Comm_create(MPI_Comm comm, MPI_Group group, MPI_Comm *new_comm)
```

New communicator based on group.

See MPI code: `groups/`

```
/* Extract the original group handle */
MPI_Comm_group(MPI_COMM_WORLD, &world_group);

/* Divide tasks into two distinct groups based upon rank */
int mygroup = 0;
if(rank >= NPROCS/2) {
    mygroup = 1;
}

int ranks1[4]= {0,1,2,3}, ranks2[4]= {4,5,6,7};
/* These arrays specify the rank to be used
 * to create 2 separate process groups.
 */
MPI_Group_incl(world_group, NPROCS/2, ranks1, &sub_group[0]);
MPI_Group_incl(world_group, NPROCS/2, ranks2, &sub_group[1]);

/* Create new new communicator and then perform collective communications */
MPI_Comm_create(MPI_COMM_WORLD, sub_group[mygroup], &sub_group_comm);
// Summing up the value of the rank for all processes in my group
MPI_Allreduce(&sendbuf, &recvbuf, 1, MPI_INT, MPI_SUM, sub_group_comm);

MPI_Group_rank(sub_group[mygroup], &group_rank);
printf("Rank= %d; Group rank= %d; recvbuf= %d\n",rank,group_rank,recvbuf);
```

The background of the slide features a large, light gray watermark of the Stanford University seal. The seal is circular and contains a redwood tree in the center, with the text "LELAND STANFORD JUNIOR UNIVERSITY" around the top and "DIE LUFT DER FREIHEIT WEHT" around the bottom. The year "1891" is at the very bottom, and there are stars on the sides.

PROCESS TOPOLOGIES

Process topologies

- Many problems are naturally mapped to certain topologies such as **grids**.
- This is the case for example for matrices, or for 2D and 3D structured grids.
- The two main types of topologies supported by MPI are **Cartesian grids and graphs**.
- MPI topologies allow simplifying many common MPI tasks.
- MPI topologies are virtual — there may be no relation between the physical structure of the network and the process topology.

Advantages of using topologies

- **Convenience:** virtual topologies may be useful for applications with specific communication patterns.
- **Communication efficiency:** a particular implementation may optimize the process mapping based upon the physical characteristics of a given parallel machine.
 - For example nodes that are nearby on the grid (East/West/North/South neighbors) may be close in the network (lowest communication time).
- The mapping of processes onto an MPI virtual topology is dependent upon the MPI implementation.

MPI functions for topologies

Many functions are available.

We only cover the basic ones.

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims,  
    int *dims, int *periods, int reorder,  
    MPI_Comm *comm_cart)
```

`ndims` number of dimensions

`dims[i]` size of grid along dimension `i`. Should not exceed the number of processes in `comm_old`.

The array `periods` is used to specify whether or not the topology has wraparound connections. If `periods[i]` is non-zero, then the topology has wraparound connections along dimension `i`.

`reorder` is used to determine if the processes in the new group are to be reordered or not. If `reorder` is false, then the rank of each process in the new group is identical to its rank in the old group.

Example

0 (0,0)	1 (0,1)
2 (1,0)	3 (1,1)
4 (2,0)	5 (2,1)

The processes are ordered according to their rank row-wise in increasing order.

Periodic Cartesian grids

```
int ndims = 2; // 3x2 2D grid
int dims[2];
dims[0] = 3; // rows
dims[1] = 2; // columns
assert(nprocs >= dims[0]*dims[1]);
int periods[2]; periods[0] = 1; periods[1] = 1;
int reorder = 1;
MPI_Cart_create(MPI_COMM_WORLD, ndims, dims,
                periods, reorder, &comm_cart);
```

- **We chose periodicity along the first dimension (`periods[0]=1`)** which means that any reference beyond the first or last entry of any row will be wrapped around cyclically.
- For example, row index `i=-1` is mapped into `i=2`.
- **There is no periodicity imposed on the second dimension.** Any reference to a column index outside of its defined range results in an error. Try it!

Obtaining your rank and coordinates

```
int MPI_Cart_rank(MPI_Comm comm_cart,  
    int *coords, int *rank)  
int MPI_Cart_coords(MPI_Comm comm_cart, int rank,  
    int maxdims, int *coords)
```

- This allows retrieving a rank or the coordinates in the grid. This may be useful to get information about other processes.
- `coords` are the Cartesian coordinates of a process.
- Its size is the number of dimensions.
- Remember that the function `MPI_Comm_rank` is still available to query your own rank.
- See MPI code: `mpi_cart/`

```
/* Get my rank in the new topology */
```

```
int my2drank;
```

```
MPI_Comm_rank(comm_cart, &my2drank);
```

```
/* Get my coordinates */
```

```
int mycoords[2];
```

```
MPI_Cart_coords(comm_cart, my2drank, 2, mycoords);
```

```
/* Get coordinates of process below me */
```

```
int rank_down, coords[2];
```

```
coords[0] = mycoords[0]+1; // i coordinate (one row below in matrix)
```

```
coords[1] = mycoords[1];
```

```
MPI_Cart_rank(comm_cart, coords, &rank_down);
```

```
/* Get coordinates of process to my right */
```

```
int rank_right;
```

```
coords[0] = mycoords[0];
```


```
coords[1] = mycoords[1]+1; // j coordinate (to the right in matrix)
```


```
MPI_Cart_rank(comm_cart, coords, &rank_right);
```

Getting the rank of your neighbors


```
int MPI_Cart_shift(MPI_Comm comm_cart, int dir,  
    int s_step, int *rank_source, int *rank_dest)
```

- `dir` direction
- `s_step` length shift
- `rank_dest` contains the group rank of the neighboring process in the specified dimension and distance.
- `rank_source` is the rank of the process for which the calling process is the neighboring process in the specified dimension and distance.
- Thus, the group ranks returned in `rank_dest` and `rank_source` can be used as parameters for `MPI_Sendrecv()`.





MPI virtual topology reflects the
topology of the underlying hardware
network



Yes

No



Start the presentation to activate live content

If you see this message in presentation mode, install the add-in or get help at PollEv.com/app



Total Results: 0

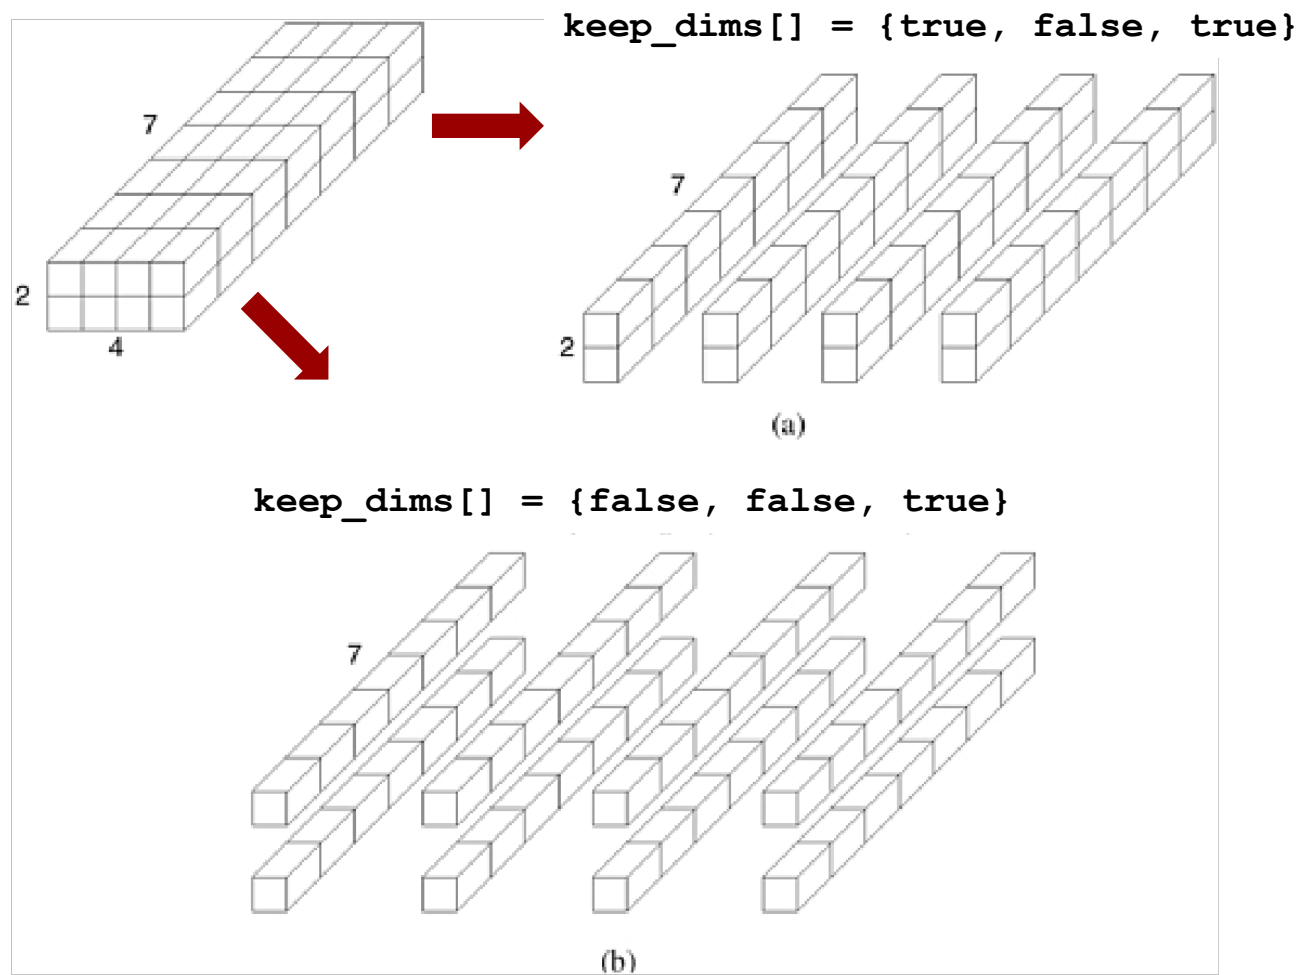

Splitting a Cartesian topology

- It is very common that one wants to split a Cartesian topology along certain dimensions.
- For example, we may want to create a group for the columns or rows of a matrix.


```
int MPI_Cart_sub(MPI_Comm comm_cart,  
                int *keep_dims, MPI_Comm *comm_subcart)
```

- `keep_dims` boolean flag that determines whether that dimension is retained in the new communicators or split, e.g., if false then a split occurs.


Example




Application example: 2D partitioning



First column
contains b



Send b to the
diagonal
processes



Send b down
each column.
Broadcast!

Start with 2D
communicator

Use column group

2D topology for matrix

```
int periods[2];
periods[ROW] = periods[COL] = 1;
/* We will use wrap-around connections. */

MPI_Comm comm_2d;
MPI_Cart_create(MPI_COMM_WORLD, 2, dims, periods, 1, &comm_2d);

/* Get my rank in the new topology */
int my2drank;
MPI_Comm_rank(comm_2d, &my2drank);

/* Get my coordinates */
int mycoords[2];
MPI_Cart_coords(comm_2d, my2drank, 2, mycoords);
```


Send to diagonal block

```
// Send to diagonal block
if(mycoords[COL] == 0 && mycoords[ROW] != 0) {
    /* I'm in the first column */
    int drank;
    int coords[2];
    coords[ROW] = mycoords[ROW];
    coords[COL] = mycoords[ROW]; // coordinates of diagonal block
    MPI_Cart_rank(comm_2d, coords, &drank); // 2D communicator
    /* Send data to the diagonal block */
    MPI_Send(&b[0], nlocal, MPI_FLOAT, drank, 1, comm_2d);
}


// Receive from column 0
if(mycoords[ROW] == mycoords[COL] && mycoords[ROW] != 0) {
    /* I am a diagonal block */
    int col0rank;
    int coords[2];
    coords[ROW] = mycoords[ROW];
    coords[COL] = 0; // Receiving from column 0
    MPI_Cart_rank(comm_2d, coords, &col0rank); // 2D communicator
    MPI_Recv(&b[0], nlocal, MPI_FLOAT, col0rank, 1, comm_2d,
            MPI_STATUS_IGNORE);
}
```

Column-wise broadcast

```
/* Create the column-based sub-topology */
MPI_Comm comm_col;
int keep_dims[2];
keep_dims[ROW] = 1;
keep_dims[COL] = 0;
MPI_Cart_sub(comm_2d, keep_dims, &comm_col);

/* Broadcast inside column */
int drank;
int coord = mycoords[COL]; // Coordinate in 1D column topology
MPI_Cart_rank(comm_col, &coord, &drank);
MPI_Bcast(&b[0], nlocal, MPI_FLOAT, drank, comm_col);
```

matvec2D



Reduction!
Use row group


See MPI code: `matvec2D/`

Code for row reduction

```
/* Create the row-based sub-topology */
MPI_Comm comm_row;
int keep_dims[2];
keep_dims[ROW] = 0;
keep_dims[COL] = 1;
MPI_Cart_sub(comm_2d, keep_dims, &comm_row);

// Row-wise reduction
int col0rank;
int coord = 0; // Coordinate in 1D row topology
MPI_Cart_rank(comm_row, &coord, &col0rank);
MPI_Reduce(&px[0], &x[0], nlocal, MPI_FLOAT, MPI_SUM, col0rank, comm_row);
```

Topologies for finite-element calculations



- A typical situation is that processes need to communicate with their neighbors.
- This becomes complicated to organize for **unstructured grids**.
- In that case, **graph topologies** are very convenient. They allow defining a neighbor relationship in a general way, using a graph. Example: `MPI_Graph_create`
- Examples of collective communications:
 - `MPI_neighbor_allgather()` : gather data, and all processes get the result
 - `MPI_neighbor_alltoall()` : processes send to and receive from all neighbor processes